

Confluent Stream Processing using Apache Kafka™ Streams & ksqlDB

Exercise Book

Version 7.0.0-v1.0.1

Table of Contents

Copyright & Trademarks	1
Lab 00 Introduction	2
a. Introduction	2
b. Continued Learning After Class	6
Lab 01 Introduction to Kafka Streams	8
a. Scaling a Kafka Streams Application	8
Lab 02 Working with Kafka Streams	18
a. Anatomy of a Kafka Streams App	18
b. Working with JSON	27
Lab 03 Introduction to ksqlDB	32
a. Introduction to ksqlDB	32
b. Using the ksqlDB REST API	40
c. Creating connectors with ksqlDB	46
Lab 04 Using ksqlDB	57
a. Using ksqlDB	57
Lab 06 Windowing & Aggregations	65
a. Windowing & Aggregations	65
Lab 07 Joins	73
a. Joining Two Streams	73
Lab 08 Custom Processing	82
a. Using the Processor API	82
Lab 09 Testing, Monitoring and Troubleshooting	87
a. Building Unit Tests	87
b. Integration Tests using Embedded Kafka	93
c. Getting Metrics from a Kafka Streams Application	97
d. Using JConsole to monitor a Streams App	101
e. Monitoring a Kafka Streams App in Confluent Control Center	107
Lab 11 Security	112
a. Securing a Kafka Streams Application	112
Appendix A: Running All Labs with Docker	121
Running Labs in Docker for Desktop	121
Running the Exercise Applications	122

Copyright & Trademarks

Copyright © Confluent, Inc. 2014-2022. [Privacy Policy](#) | [Terms & Conditions](#).

Apache, Apache Kafka, Kafka and the Kafka logo are trademarks of the
[Apache Software Foundation](#)

Lab 00 Introduction

a. Introduction

This document provides Hands-On Exercises for the course Confluent Stream Processing using Apache Kafka™ Streams & ksqlDB. You will use a setup that includes a virtual machine (VM) configured with Apache Kafka and Confluent tools to manage your data and clusters.

Alternative Lab Environment

As an alternative you can:

- ¥ Download the VM to your laptop and run it in VirtualBox. Make sure you have the newest version of VirtualBox installed. Download the VM from this link:

<https://s3.amazonaws.com/confluent-training-images-us-east-1/training-ubuntu-20-04-jan2022.ova>

- ¥ If you have installed Docker for Desktop on your Mac or Windows 10 Pro machine then you can run the labs there. But please note that your trainer might not be able to troubleshoot any potential problems if you are running the labs locally. If you choose to do this, follow the instructions at ! [Running Labs in Docker for Desktop](#).

Command Line Examples

Most exercises contain commands that must be run from the command line. These commands will look like this:

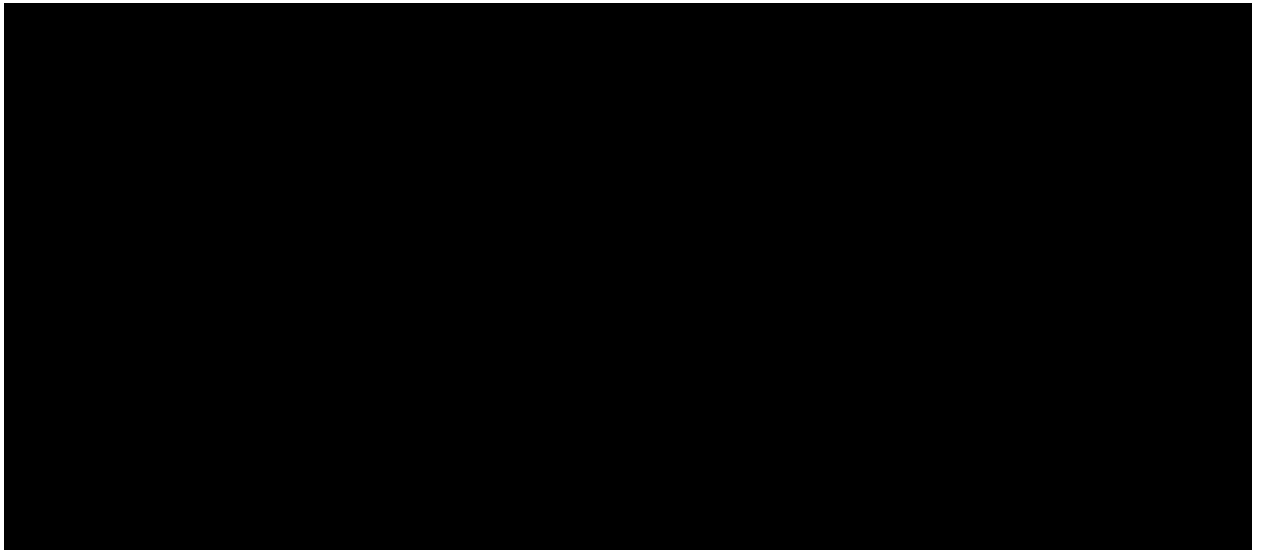
```
$ pwd
/home/trai ni ng
```

Commands you should type are shown in bold; non-bold text is an example of the output produced as a result of the command.

Preparing the Labs

Welcome to your lab environment! You are connected as user training, password training.

If you haven't already done so, you should open the Exercise Guide that is located on the lab virtual machine. To do so, open the Confluent Training Exercises folder that is located on the lab virtual machine desktop. Then double-click the shortcut that is in the folder to open the Exercise Guide.



Copy and paste works best if you copy from the Exercise Guide on your lab virtual machine.

¥ Standard Ubuntu keyboard shortcuts will work: **Ctrl +C** ! Copy, **Ctrl +V** ! Paste

¥ In a Terminal window: **Ctrl +Shift+C** ! Copy, **Ctrl +Shift+V** ! Paste.

If you find these keyboard shortcuts are not working you can use the right-click context menu for copy and paste.

1. Open a terminal window
2. Clone the source code repository to the folder **confluent-streams** in your home directory:

```
$ cd ~
$ git clone --depth 1 --branch 7.0.0-v1.0.1 \
  https://github.com/confluentinc/training-ksql-and-streams-src.git
$ cd confluent-streams
```

”

If you chose to select another folder for the labs then note that many of our samples assume that the lab folder is `~/confluent-streams`. You will have to adjust all those command to fit your specific environment.

3. Navigate to the `confluent-streams` folder:

```
$ cd ~/confluent-streams
```

4. Start the Kafka cluster:

```
$ docker-compose up -d zookeeper kafka
```

You should see something similar to this:

```
Creating network "confluent-streams_kafka-net" with the default
driver
Creating kafka      ... done
Creating zookeeper  ... done
```

In the first steps of each exercise, you may launch the containers needed for the exercise with a `docker-compose up` command. Simply typing `docker-compose up -d` will start all of the containers defined in the `docker-compose.yml` file. You can start fewer containers by specifying only those you want to run, for example: `docker-compose up -d zookeeper kafka`.

The majority of the exercises use the `docker-compose.yml` file in the `~/confluent-streams` directory. The `docker-compose up` command will search up the directory hierarchy until it finds a `docker-compose.yml` file, so the one in the `confluent-streams` directory will usually be used. The exception is the `docker-compose.yml` file used in the security exercise as this has additional security settings. See the comments at the beginning of Lab 11 Securing a Kafka Streams Application.

If at any time you want to get your environment back to a clean state use `docker-compose down` to end all of your containers. Then return to your last `docker-compose up` to get back to the beginning of an exercise.

Exercises do not need to be completed in order. You can start from the beginning of any exercise at any time.

If you want to completely clear out your docker environment use the script on the VM: `docker-nuke.sh`. The nuke script will forcefully end all of your running docker containers.

5. Monitor the cluster with:

```
$ docker-compose ps
```

Name	Command	State	Ports
kafka	/etc/confluent/docker/run	Up	0.0.0.0: 9092->9092/tcp
zookeeper	/etc/confluent/docker/run	Up	2181/tcp, 2888/tcp, 3888/tcp

All services should have `State` equal to `Up`.

6. You can also observe the stats of Docker on your VM:

```
$ docker stats
```

CONTAINER ID	NAME	CPU %	MEM USAGE / LIMIT	MEM %	NET I/O
ab9c97077e94	zookeeper	0.14%	88.14MiB / 9.737GiB	0.88%	106kB / 130kB
ff47bece9e4f	kafka	1.17%	421.8MiB / 9.737GiB	4.23%	646kB / 522kB

Cleanup

1. Press **Ctrl +C** to exit the Docker statistics.
2. Shut down your Kafka cluster with the **docker-compose down -v** command.

b. Continued Learning After Class

Once the course ends, the VM in Content Raven will terminate and you will no longer have access to it. However, you can still download the VM onto your own machine or use Docker locally to revisit these materials. We encourage you to bring up your own test environment, explore configuration files, inspect scripts, and perform tests. Here are some activities we encourage to reinforce your learning:

- ¥ Revisit the exercises in this manual
- ¥ Summarize and discuss the student handbook with your peers
- ¥ Consult the README in this public repository for more resources and your own development playground: <https://github.com/confluentinc/training-ksql-and-streams-src>

Conclusion

In this lab you have prepared and tested the Lab Environment. Finally you have created your Apache Kafka cluster that will be used in subsequent exercises.

STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 01 Introduction to Kafka Streams

a. Scaling a Kafka Streams Application

In this exercise, we're going to write a Kafka producer in either Python or Java that generates a stream of temperature readings for a set of weather stations. We are also writing a simple Kafka Streams application that will consume this topic and calculate the maximum temperature per station per time window. We will then run this application in a single instance and later scale it up to several instances. We will monitor the throughput with the Confluent Control Center.

Prerequisites

Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)

Running the Kafka Cluster and Confluent Control Center

To be able to run Kafka Streams applications we need a working Kafka cluster. We will run one consisting of a single broker and zookeeper.

1. Locate the file `docker-compose.yml` in the `~/confluent-streams` directory.



This `docker-compose.yml` file will be used to run a simple Kafka cluster as a backend for our Kafka Streams applications. If you're curious, please open the file and analyze its contents.

2. Navigate to the folder `labs/scaling`:

```
$ cd ~/confluent-streams/labs/scaling
```

3. Run the Kafka cluster:

```
$ docker-compose up -d zookeeper kafka ksqldb-server control-center
```

Wait a couple of minutes until the cluster is initialized.

4. Create the two topics `temperature-readings` and `max-temperatures`, each with 3 partitions using these commands:

```
$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 3 \
  --topic temperature-readings

$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 3 \
  --topic max-temperatures
```

Creating the Producer

Now it is time to create a temperature readings producer. You can do it either in Python or Java. We start with the Python producer. If you prefer Java then move ahead to ! [here](#).

Create the Producer in Python

1. Navigate to the folder `labs/scaling`, and launch VS Code:

```
$ cd ~/confluent-streams/labs/scaling/temp-producer
$ code .
```

#

If a pop-up appears when VS Code opens, check the box and click `Yes, I trust the authors`.

2. Locate the file `main.py`.
3. Inspect the code.

- ! We are defining a few (temperature measurement) stations, their respective average temperatures and last measured temperatures
- ! we're using the Confluent Python client for Kafka to create a producer
- ! every ~100 ms we're generating a temperature reading for one of the randomly selected station.

4. Return to the terminal window, and install the python prerequisite. Note this may already be installed from a previous exercise:

```
$ pip3 install --upgrade pip
$ pip3 install confluent-kafka
```

5. Run the producer:

```
$ python3 main.py
```

6. In another terminal window run the `kafka-console-consumer` to display the `temperature-readings` topic:

```
$ kafka-console-consumer \
  --bootstrap-server kafka:9092 \
  --from-beginning \
  --max-messages 25 \
  --topic temperature-readings \
  --property print.key=true \
  --property key.separator=", "
```

You should see an output similar to this:

```
S-06, {"station": "S-06", "temperature": -1}
S-03, {"station": "S-03", "temperature": 8}
S-03, {"station": "S-03", "temperature": 9}
S-06, {"station": "S-06", "temperature": 0}
S-08, {"station": "S-08", "temperature": 31}
S-09, {"station": "S-09", "temperature": -7}
...
```

The `key` is the station and the `value` is a JSON object with the station and the temperature in degree Celsius.

Create the Producer in Java

1. Navigate to the folder `labs/scaling`, and launch VS Code:

```
$ cd ~/confluent-streams/labs/scaling/temp-producer
$ code .
```

2. Locate the file `build.gradle` and analyze its content. It is the build file for a simple Kafka client.
3. Locate the file `TempProducer.java` in the subfolder `src/main/java/streams` and open it. Analyze the file and make sure you understand the code. If necessary discuss with your peers.
4. Notice the file `log4j.properties` in the folder `src/main/resources` that configures logging for the producer.
5. Use Run ! Start Debugging in VS Code or `./gradlew run` in the terminal to run the Java producer.



The first time you run the debugger it may take extra time while resources are downloaded.

You should see an output similar to this:

```
The record is: S-06, {"station": "S-06", "temperature": -1}
The record is: S-03, {"station": "S-03", "temperature": 8}
The record is: S-03, {"station": "S-03", "temperature": 9}
The record is: S-06, {"station": "S-06", "temperature": 0}
The record is: S-08, {"station": "S-08", "temperature": 31}
The record is: S-09, {"station": "S-09", "temperature": -7}
...
```

The **key** is the station and the **value** is a JSON object with the station and temperature in degree Celsius.

6. From another terminal window, run the `kafka-console-consumer` to display the `temperature-readings` topic:

```
$ kafka-consol e-consumer \  
E --bootstrap-server kafka:9092 \  
E --from-begi nni ng \  
E --max-messages 25 \  
E --topi c temperature-readi ngs \  
E --property print.key=true \  
E --property key.separator=", "
```

You should see an output similar to this:

```
S-06, {"station": "S-06", "temperature": -1}  
S-03, {"station": "S-03", "temperature": 8}  
S-03, {"station": "S-03", "temperature": 9}  
S-06, {"station": "S-06", "temperature": 0}  
S-08, {"station": "S-08", "temperature": 31}  
S-09, {"station": "S-09", "temperature": -7}  
...
```

7. Jump to the next section "Writing the Kafka Streams Application".

Writing the Kafka Streams Application

1. Open a terminal window and navigate to the `streams-app` folder:

```
$ cd ~/confluent-streams/abs/scaling/streams-app
```

2. In the folder `confluent-streams/abs/scaling/streams-app` locate the file `build.gradle` and analyze its content.



In addition to the usual libraries we also load the `monitoring-interceptors` library to be able to integrate with the Confluent Control Center and the `kafka-json-serializer` library for the JSON serde.

3. We will be using the Kafka Streams application called `StreamsApp.java` in the subfolder `src/main/java/streams` for this exercise.
4. You may choose to launch VS Code with `code .` to build and run the application. Or simply use gradle with `./gradlew run`.
5. Run an instance of `kafka-consol e-consumer` to display the `max-temperatures` topic.

Note it may take some time for max temperatures to appear:

```
$ kafka-console-consumer \
  --bootstrap-server kafka:9092 \
  --from-beginning \
  --topic max-temperatures
```

!

Is the output surprising to you? Why? It is because of the nature of the `commit.interval.ms` property and how it relates to the output of KTables to Kafka topics. This will be discussed later in the course.

Ignore the warnings:

```
2022-05-11 21:45:52 WARN ConsumerConfig:362 - The
configuration 'admin.retry.backoff.ms' was supplied but isn't
a known config. 2022-05-11 21:45:52 WARN ConsumerConfig:362 -
The configuration 'admin.retries' was supplied but isn't a
known config.
```

6. Open Confluent Control Center at <http://localhost:9021>
7. In Control Center, click on `CONTROLCENTER.CLUSTER` and then under Consumers, monitor the consumer lag for the `streams-app-v.0.1.0`. Note that the consumer group falls more and more behind:

- Alternatively, you can use the Kafka tool `kafka-consumer-groups` to check the consumer lag:

```
$ kafka-consumer-groups
Ê --bootstrap-server kafka:9092
Ê --group streams-app-v0.1.0
Ê --describe
GROUP                                TOPIC                                PARTITION  CURRENT-OFFSET
LOG-END-OFFSET  LAG    CONSUMER-ID...
streams-app-v0.1.0 temperature-readings 0          551
2417            1866    streams-app...
streams-app-v0.1.0 temperature-readings 1          363
5395            5032    streams-app...
streams-app-v0.1.0 temperature-readings 2          580
3614            3034    streams-app...
```

Scaling the Kafka Streams Application

- Open another terminal window to run another instance of your Kafka Streams application:

```
$ cd ~/confluent-streams/labs/scaling/streams-app
$ ./gradlew run
```

- In Confluent Control Center observe how the throughput of the streams app nearly doubles.
- Note that we have now two consumer instances listed (recognizable by their ID):

4. Also observe that the consumer lag increases more slowly
5. Now scale the streams app to 3 instances and again monitor an increase in throughput and reduction in consumer lag.
6. Finally scale the app again, this time to 4 instances. Monitor the throughput after scaling the app. What are you observing? Explain your observation. See the conclusion for an explanation of what happens here.

Optional: Using the ksqlDB CLI

We can achieve the same results using ksqlDB!

1. Use the ksqlDB CLI to play with the data:

```
$ ksql http://ksql-db-server:8088
```

2. Set the starting point of your queries to `earliest`:

```
ksql > SET 'auto.offset.reset' = 'earliest';
```

3. Create a stream for the source topic:

```
ksql > CREATE STREAM temperatures(station STRING, temperature INTEGER)
Ê      WITH(KAFKA_TOPIC='temperature-readings',
Ê      VALUE_FORMAT='JSON');
```

4. Create a table that shows the maximum, minimum, and average temperatures per station per minute:

```
ksql > CREATE TABLE temp_agg_per_min AS
Ê      SELECT station,
Ê      max(temperature) AS max,
Ê      min(temperature) AS min,
Ê      sum(temperature) / count( * ) AS avg
Ê      FROM temperatures
Ê      WINDOW TUMBLING (SIZE 1 MINUTE)
Ê      GROUP BY station;
```

5. Inspect the aggregated temperature data as new records flow in from the producer.

```
ksql > SELECT station, max, min, avg FROM temp_agg_per_min EMIT
CHANGES;
```

Press **CTL-C** to terminate the query.

6. Exit ksqldb with **Ctrl +D**.

Cleanup

1. Stop the producer, consumers, and stream application with **Ctrl +C** in the terminal or Run ! Stop Debugging in VScode.
2. Shut down your Kafka cluster with the **docker-compose down -v** command.

Conclusion

In this exercise we have created a Kafka Streams application that processed an input topic and produced an output topic. First we ran only one application instance and then we scaled the application up to several instances. We noticed a significant boost in throughput until the number of instances was greater than the number of partitions of the input topic. At this point the additional application instances were sitting there idle.



In the solutions folder, the Java producer and Kafka Streams app include Dockerfiles so that they can be deployed as containers. A separate `docker-compose. services. yml` has also been provided to start the microservices.

STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 02 Working with Kafka Streams

This lab contains 2 exercises:

¥ Anatomy of a Kafka Streams App

¥ Working with JSON

a. Anatomy of a Kafka Streams App

In this exercise, you will create a Kafka Streams application and deploy it using Gradle (or, optionally, Maven). The purpose of this exercise is to illustrate the structure of Kafka Streams application code and the routine of deploying code with build tools.

The application itself reads data from a topic whose keys are integers and whose values are sentence strings. The application transforms the input so that the strings are lower-case and output to a new topic.

Prerequisites

Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)

Preparing the Kafka Cluster

1. Navigate to the folder `~/confluent-streams/labs/working-streams`:

```
$ cd ~/confluent-streams/labs/working-streams
```

2. From within the `working-streams` folder run the cluster with the following command:

```
$ docker-compose up -d zookeeper kafka
Creating network "confluent-streams_kafka-net" with the default
driver
Creating kafka          ... done
Creating zookeeper      ... done
```

3. Double check that the cluster is up and running:

```
$ docker-compose ps
```

you should see something similar to this:

Ports	Name	Command	State
0.0.0.0:9092->9092/tcp	kafka	/etc/confluent/docker/run	Up
2181/tcp, 2888/tcp, 3888/tcp	zookeeper	/etc/confluent/docker/run	Up

Make sure all services have **State=Up**.

4. Create an input topic called **lines-topic** in Kafka:

```
$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic lines-topic
```

5. Create an output topic called **lines-lower-topic** in Kafka:

```
$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic lines-lower-topic
Created topic "lines-lower-topic".
```

6. Let's now check what topics are in Kafka using this command:

```
$ kafka-topics \
  --bootstrap-server kafka:9092 \
  --list
```

We should see something like this:

```
__confluent.support.metrics
lines-lower-topic
lines-topic
```



If you need to delete a topic, say the one with name `<topic name>`, (e.g. to start over) you can use this command:

```
$ kafka-topics \
  --bootstrap-server kafka:9092 \
  --delete \
  --topic <topic name>
```

Now we are ready to create, build and run our first Kafka Streams application. We will first build and run the application using Java in conjunction with Gradle. Optionally, we will also see how to build and run the application with Maven.

Authoring the Kafka Streams Application using Java & Gradle

1. Navigate to the folder `gradle-sample` within the `working-streams` folder, and launch VS Code:



Be certain to include the period in the `code .` command below. That indicates starting VS Code in the current directory - otherwise some references may not correctly resolve.

```
$ cd ~/confluent-streams/abs/working-streams/gradle-sample
$ code .
```

2. In this folder locate the `build.gradle` file and open it to analyze its content.



The external dependencies for a simple Kafka Streams app are the `kafka-clients` and `kafka-streams` libraries. We also add the `slf4j-log4j12` dependency for logging purposes.

3. In VS Code, open the file `MapSample.java` in the subfolder `src/main/java/streams` and inspect its content. From here, you can challenge yourself to complete the TODOs, or you can move forward to see a step-by-step walkthrough of the code. If you decide to challenge yourself, you can always peek at the corresponding subfolder in `~/confluent-streams/solutions/` if you get stuck.
4. Now we start to add actual Kafka Streams application logic. We will start with the configuration part. Please add this code snippet to the `main` method of the class (right after the initial `println`):

```
Properties settings = new Properties();
settings.put(StreamsConfig.APPLICATION_ID_CONFIG, "map-sample-
v0.1.0");
settings.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "kafka: 9092");
```

We're providing an ID to our application and tell it where to find the Kafka cluster. This is the minimal configuration needed!

5. Next we will define the topology of our Kafka Streams application. Add this snippet right after the configuration part:

```
final Serde<String> stringSerde = Serdes.String();
StreamsBuilder builder = new StreamsBuilder();
KStream<String, String> lines = builder
    .stream("lines-topic", Consumed.with(stringSerde, stringSerde));
KStream<String, String> transformed = lines
    .mapValues(value -> value.toLowerCase());
transformed.to("lines-lower-topic", Produced.with(stringSerde,
stringSerde));
Topology topology = builder.build();
```

We're defining a builder for the topology, use it to create a `KStream` from the Kafka topic `lines-topic` using `String` Serdes for both key and value of the messages. Then we're using a `mapValues` function on the `KStream` to convert the `value` into all lower case. Finally we're writing the result into the Kafka topic `lines-lower-topic`. Then we build the topology.

6. With the settings and the topology at hand we can now create the Streams app:

```
KafkaStreams streams = new KafkaStreams(topology, settings);
```

Our application will now start consuming data from the input topic, transform it and write it to the output topic.

7. To have our application terminate in an orderly way when requested without leaving any resource leaks behind we add a shutdown hook at the end of the main method:

```
final CountDownLatch latch = new CountDownLatch(1);
Runtime.getRuntime().addShutdownHook(new Thread(() -> {
    System.out.println("### Stopping Map Sample Application ###");
    streams.close();
    latch.countDown();
}));

try{
    streams.start();
    latch.await();
} catch (final Throwable e) {
    System.exit(1);
}
System.exit(0);
```

This Shutdown Hook will be executed when the application receives a `SIG_TERM` signal. The `CountDownLatch` is used as a best practice to avoid rare cases of deadlock. Notice the use of the `await()` method after the application starts.

8. Within the project folder `gradle-sample` locate the folder for the Java resources `src/main/resources`. In this folder we have a file `log4j.properties`. This file is used to configure the logger for our Kafka Streams app.



Use `INFO` instead of `WARN` for the `rootLogger` if you want to be more verbose in the logs.

And that's all we need. Our first Kafka Streams app is ready to go! Use Run ! Start Debugging to run your code.

In the VS Code DEBUG CONSOLE tab you should see something like this:

```
*** Starting Map Sample Application ***  
2018-07-04 13:47:36 WARN ConsumerConfig: 287 - The configuration  
'admin.retries' was supplied but isn't a known config.
```

At this time you can safely ignore the WARN log items.

!

At this time nothing will happen since we did not yet produce any data in the `lines-topic` in Kafka. Let your Kafka Streams app remain running in the VS Code debugger. This will be our next task after we have shown how to build the same app using Maven instead of Gradle.
If you want to skip the Maven part then go to [Producing some Input Data](#).

Optional: Build and Run the application with Java & Maven

Here we're basically showing the same steps as in the previous section, except we will now use Maven instead of Gradle.

1. Navigate to the `maven-sample` folder, and launch VS Code:

```
$ cd ~/confluent-streams/abs/working-streams/maven-sample  
$ code .
```

2. Locate the `pom.xml` file, open it and analyze its content.



this is the minimum Maven file required to compile, package and run a Java application that has external dependencies on the 2 required libraries `kafka-clients` and `kafka-streams`. It also uses the `slf4j-log4j12` library for logging.

3. Locate the file `MapSample.java` in the subfolder `src/main/java/streams` and open it. Double check that it looks similar to the one you created in the Gradle example.
4. Note the same file `log4j.properties` in the folder `src/main/resources` as in the Gradle example.



Use `INFO` instead of `WARN` for the `rootLogger` if you want to be more verbose in the logs.

5. Solution code has been provided in this exercise, so our first Kafka Streams app is ready to go! Use Run ! Start Debugging to run your code.

You should see something like this:

```
*** Starting Map Sample Application ***
2018-07-04 13:53:12 WARN ConsumerConfig: 287 - The configuration
'admin.retries' was supplied but isn't a known config.
```

At this time you can safely ignore the WARN log items.



At this time nothing will happen since we did not yet produce any data in the `lines-topic` in Kafka. Let's do this next. Let your Kafka Streams app remain running in the VS Code debugger.

Producing some Input Data

We're going to use the `kafka-console-producer` tool to create some input data in the topic `lines-topic`.

1. Open a new terminal window and navigate to the `working-streams` folder:

```
$ cd ~/confluent-streams/abs/working-streams
```

2. From a terminal window execute this command:

```
$ cat << EOF | kafka-console-producer \
  --bootstrap-server kafka:9092 \
  --property "parse.key=true" \
  --property "key.separator=: " \
  --topic lines-topic
1: "Kafka powers the Confluent Streaming Platform"
2: "Events are stored in Kafka"
3: "Confluent contributes to Kafka"
EOF
```

This writes 3 entries into the topic called `lines-topic` using `String` serializers for both key and value.

Reading the Transformed Messages

Here we're using the `kafka-console-consumer` tool to read from the output topic. We're again using `String` deserializers for key and value.

1. In the same terminal window, run:

```
$ kafka-console-consumer \
  --bootstrap-server kafka:9092 \
  --from-beginning \
  --topic lines-lower-topic
```

You should see this:

```
"kafka powers the confluent streaming platform"
"events are stored in kafka"
"confluent contributes to kafka"
```



You might need to be a bit patient until the messages appear. It can take a few seconds to a minute or so depending on the performance of your computer.

Cleanup

1. Terminate the Kafka console consumer by pressing `Ctrl +C`.
2. Terminate your Kafka Streams application with Run ! Stop Debugging. If you used the `./gradlew run` command instead, you can terminate with `Ctrl +C`.
3. Shut down your Kafka cluster with the `docker-compose down -v` command.

Conclusion

We have created our first complete Kafka Streams application. We built the application using Gradle, and then again with Maven. We have used the command line tools provided by Kafka to produce input data and display the transformed output data.

STOP HERE. THIS IS THE END OF THE EXERCISE.

b. Working with JSON

The purpose of this exercise is to learn how to create serializers and deserializers for custom Java objects.

In this case, we will create a Serde for an object that records temperature data using the `KafkaJsonSerializer` and `KafkaJsonDeserializer` helper classes. The application itself reads temperature data from an input topic, filters for temperatures higher than 25 degrees, and outputs that data to a new output topic.

Prerequisites

1. Navigate to this lab's folder `~/confluent-streams/labs/working-streams`:

```
$ cd ~/confluent-streams/labs/working-streams
```

2. If it is not already running, start the Kafka cluster:

```
$ docker-compose up -d zookeeper kafka
```

Do not proceed until all services are up and running; test with:

```
$ docker-compose ps
```

and assert that all services are in state `Up`.

3. Create an input topic called `temperatures-topic` and an output topic called `high-temperatures-topic` in Kafka:

```
$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic temperatures-topic

$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic high-temperatures-topic
```

Writing the Kafka Streams App

1. Open a new terminal window and navigate to the `working-streams/json-sample` folder, and launch VS Code:

```
$ cd ~/confluent-streams/labs/working-streams/json-sample
$ code .
```

2. Locate the file `build.gradle` in this folder and open it to analyze its content.



Compared to the `build.gradle` file in the previous exercise we have added the `kafka-json-serIALIZER` library for the JSON serializer/deserializer

3. In the subfolder `src/main/java/streams` locate the file `JsonSample.java` and familiarize yourself with the code. It basically does the configuration of the Kafka Streams app but the interesting code is missing. As before, you can challenge yourself to implement the missing code or follow the step-by-step instructions in this book.
4. To create a Serde (Serializer/Deserializer) for JSON formatted data, add this code to the function `getJsonSerde()` after the TODO comment:

```

Map<String, Object> serdeProps = new HashMap<>();
serdeProps.put("json.value.type", TempReading.class);

final Serializer<TempReading> temperatureSerializer = new
KafkaJsonSerializer<>();
temperatureSerializer.configure(serdeProps, false);

final Deserializer<TempReading> temperatureDeserializer = new
KafkaJsonDeserializer<>();
temperatureDeserializer.configure(serdeProps, false);

return Serdes.serdeFrom(temperatureSerializer,
temperatureDeserializer);

```

We're basically using the two helper classes `KafkaJsonSerializer` and `KafkaJsonDeserializer` to create a serializer and a deserializer which in turn we then use to create a Serde.

We will use this Serde to serialize and deserialize our `TempReading` POJO.

5. The final thing to do is to define the Topology for our application. We want to keep it simple and just filter the input topic `temperatures-topic` for high temperatures (>25 degrees) and output the result to the output topic `high-temperatures-topic`. Add this code to the `getTopology()` function after the TODO comment:

```

builder.stream("temperatures-topic", Consumed.with(stringSerde,
temperatureSerde))
    .filter((key, value) -> value.temperature > 25)
    .to("high-temperatures-topic", Produced.with(stringSerde,
temperatureSerde));
return builder.build();

```



Note how the filter function uses the `value` which is an object of type `TempReading`.

6. Note the file `log4j.properties` in the folder `src/main/resources` which is used to configure logging for our application.
7. Use Run ! Start Debugging to run your code. Let your Kafka Streams app remain running in the VS Code debugger.

You should get this output:

*** Starting JSON Sample Application ***

...

Creating input Data

1. Switch back to the terminal.
2. Use the following command to generate some temperature readings in JSON format:

```
$ cat << EOF | kafka-console-producer \
Ê --bootstrap-server kafka:9092 \
Ê --property "parse.key=true" \
Ê --property "key.separator=: " \
Ê --topic temperatures-topic
"S1":{"station":"S1", "temperature": 10.2, "timestamp": 1}
"S1":{"station":"S1", "temperature": 11.2, "timestamp": 2}
"S1":{"station":"S1", "temperature": 11.1, "timestamp": 3}
"S1":{"station":"S1", "temperature": 12.5, "timestamp": 4}
"S2":{"station":"S2", "temperature": 15.2, "timestamp": 1}
"S2":{"station":"S2", "temperature": 21.7, "timestamp": 2}
"S2":{"station":"S2", "temperature": 25.1, "timestamp": 3}
"S2":{"station":"S2", "temperature": 27.8, "timestamp": 4}
EOF
```



Run this command repeatedly to generate more messages

Reading the Output

1. Use the following command to read the output generated:

```
$ kafka-console-consumer \
Ê --bootstrap-server kafka:9092 \
Ê --from-beginning \
Ê --topic high-temperatures-topic
```

You should get this output showing only readings with temperature higher than 25 degrees:

```
{"station": "S2", "temperature": 25.1, "timestamp": 3}
{"station": "S2", "temperature": 27.8, "timestamp": 4}
```


Cleanup

1. Terminate the Kafka console consumer by pressing `Ctrl +C`.
2. Terminate your Kafka Streams application with Run ! Stop Debugging.
3. Shut down your Kafka cluster with the `docker-compose down -v` command.

Conclusion

In this sample we have built a Kafka Streams application that uses custom serializer and deserializer to work with data that is JSON formatted.

STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 03 Introduction to ksqlDB

a. Introduction to ksqlDB

In this lab exercise, you will use the ksqlDB CLI to slice and dice data that is being generated in real time. The purpose is to get a sense for the streaming applications that are possible using only a SQL-like syntax. We will query running streams, apply filters and maps, and create new streams and tables derived from existing streams.

For more details on running ksqlDB in Docker containers, please see [Install ksqlDB with Docker](#).

Prerequisites

Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)

Preparing the Platform

1. Open a terminal window and navigate to the folder `~/confluent-streams`:

```
$ cd ~/confluent-streams
```

2. This exercise will use the `docker-compose.yml` file in `~/confluent-streams`, but this time we need to also run the `ksqldb-server` container.

Optional: Open the file in your editor and analyze its content. You do not necessarily need to understand all of it at this point.

3. Run the application with:

```
$ docker-compose up -d zookeeper kafka ksqldb-server
```

4. Wait until the app is up and running, that is all services are marked as `up`:

```
$ docker-compose ps
```

you should see something like this:

Name	Command	State	Ports
kafka	/etc/confluent/docker/run	Up	0.0.0.0: 9092-
ksql db-server	/etc/confluent/docker/run	Up	0.0.0.0: 8088-
zookeeper	/etc/confluent/docker/run	Up	2181/tcp, 2888/tcp, 3888/tcp

¶ You can see that we have 3 containers running on our system:

- ! the first in the list is running a Kafka broker
- ! the second container runs an instance of ksqlDB Server
- ! and the last one runs an instance of Zookeeper

All these containers run on a software defined network (SDN) called `confluent-streams_kafka-net`.

Running the ksqlDB CLI

We want now to use the ksqlDB CLI and connect it with our ksqlDB server.

5. Start the ksqlDB CLI

```
$ ksql http://ksql db-server: 8088
```

You will be greeted by the following screen:

OpenJDK 64-Bit Server VM warning: Option UseConcMarkSweepGC was deprecated in version 9.0 and will likely be removed in a future release.

[illegible]

Copyright 2017-2021 Confluent Inc.

```
CLI v7.0.0, Server v7.0.0 located at http://localhost:8088
Server Status: RUNNING
```

Having trouble? Type 'help' (case-insensitive) for a rundown of how things work!

ksql >

Working with ksqIDB CLI

Here we're going to experiment with various features of ksqlDB. We use the two topics `pageviews` and `users`.

First we will use the ksql-datagen tool to create data for us. This tool has a number of predefined data types - you see us requesting them here with the parameter `qui ckstart=users` and `qui ckstart=pagevi ews` in the commands below.

For more info on this tool, see: <https://docs.ksqldb.io/en/latest/developer-guide/test-and-debug/generate-custom-test-data/>.

6. First, create the two topics `users` and `pageviews`. Run:

```
$ kafka-topics --bootstrap-server kafka:9092 --create --partitions 1
--replication-factor 1 --topic users

$ kafka-topics --bootstrap-server kafka:9092 --create --partitions 1
--replication-factor 1 --topic pageviews
```

7. Open two more terminals and run each of the ksql-datagen commands in one of them.

```
$ ksql -datagen quickstart=users format=json topic=users \
E bootstrap-server=kafka: 9092

$ ksql -datagen quickstart=pageviews format=delimited \
E topic=pageviews bootstrap-server=kafka: 9092
```

8. Now, return to the ksqlDB CLI. To be able to work with data from Kafka we need to create either a Stream or a Table in ksqlDB. Let's create a stream from the topic `pageviews` using KSQL:

```
ksql > CREATE STREAM pageviews_original (
E viewtime bigint,
E user_id varchar,
E page_id varchar
E ) WITH (kafka_topic='pageviews', value_format='DELIMITED');
```

This command creates a Stream called `pageviews_original` from the Kafka topic `pageviews`, whose record values are encoded in CSV (here called `DELIMITED`). The three values in each record value are interpreted as fields `viewtime`, `user_id` and `page_id`.

The ksqlDB editor should answer with

```
Message
-----
Stream created
-----
```

9. We can now describe the stream:

```
ksql > DESCRIBE pageviews_original ;
```

giving us this output:

Name : PAGEVIEW_ORIGINAL

Field	Type
-------	------

VIEWTIME	BIGINT
USERID	VARCHAR(STRING)
PAGEID	VARCHAR(STRING)

For runtime statistics and query details run: DESCRIBE EXTENDED <Stream, Table>;

10. Let's create a table from the `users` topic:

```
ksql > CREATE TABLE users_original (  
    registertime BIGINT,  
    gender VARCHAR,  
    regionid VARCHAR,  
    userid VARCHAR PRIMARY KEY  
    ) WITH  
    (kafka_topic='users', value_format='JSON');
```

This command creates a Table from the Kafka topic `users`, whose records have values encoded in `JSON`. Since this is a table, we need a key by which the records from the source topic are grouped. In our case this is the field `userid`. The four fields `registertime`, `gender`, `regionid`, and `userid` should be part of the `JSON` value of the records in the topic.

11. Use `DESCRIBE users_original ;` to get a description of the table.

```
ksql > DESCRIBE users_original ;
```

giving us this output:

Name : USERS_ORIGINAL

Field	Type
-------	------

REGISTERTIME	BIGINT
GENDER	VARCHAR(STRING)
REGIONID	VARCHAR(STRING)
USERID	VARCHAR(STRING) (primary key)

For runtime statistics and query details run: DESCRIBE EXTENDED <Stream, Table>;

12. Now use `SHOW STREAMS;` and `SHOW TABLES;` to view the list of streams and tables defined in the system.

```
ksql > SHOW STREAMS;
ksql > SHOW TABLES;
```

! All ksqlDB commands need to be terminated with a semi-colon (;). ksqlDB SQL keywords such as CREATE or DESCRIBE are not case sensitive.

Querying Streams and Tables

13. Let's get some data from the `pageviews_original` stream:

! The default for where to start when selecting data is latest. If data is not being continuously loaded into your stream, this may result in your SELECT not displaying any data for some time. Set the default to be earliest in your ksqlDB CLI session with the set command: `set 'auto.offset.reset'='earliest';`

```
ksql > SELECT * FROM pageviews_original EMIT CHANGES LIMIT 10;
```

Please note the `LIMIT` clause which limits the output to 10 records. The output should look similar to this:

VIEWTIME	USERID	PAGEID
1603202197108	User_5	Page_22
1603202197145	User_6	Page_55
1603202197145	User_1	Page_82
1603202197145	User_4	Page_86
1603202197145	User_2	Page_58
1603202197145	User_8	Page_97
1603202197145	User_6	Page_88
1603202197145	User_4	Page_21
1603202197145	User_5	Page_65
1603202197145	User_5	Page_34

Limit Reached
Query terminated



If your data is displayed in columns that are too wide, you can change the column width in your ksqlDB CLI session using the set command. For example: SET CLI COLUMN-WIDTH 15;

14. Let's run the same query but this time without the **LIMIT** clause:

```
ksql > SELECT * FROM pageviews_original EMIT CHANGES;
```

you will notice that the query does not stop and continues indefinitely. This is of course expected since a stream never ends.

Hit **Ctrl +C** to end the query.

15. Now try the same with the table:

```
ksql > SELECT * FROM users_original EMIT CHANGES LIMIT 5;
```

giving us this output:

```
+-----+-----+-----+
+-----+
|USER ID          |REGISTRATION TIME    |GENDER
|REGISTRATION ID  |
+-----+-----+-----+
+-----+
|User_2           |1512186975750        |FEMALE
|Registration_4    |
|User_2           |1507817615345        |MALE
|Registration_3    |
|User_2           |1493582927082        |FEMALE
|Registration_3    |
|User_5           |1515794322305        |FEMALE
|Registration_7    |
|User_2           |1514239675179        |FEMALE
|Registration_9    |
Limit Reached
Query terminated
```




Although tables are compacted topics and can be compared to classical database tables a query on them never ends analogous to the stream.

Filtering and Mapping Operations

Similar to what we're used to from SQL we can filter data from a stream (or from a table).

16. Let's only display records from the stream whose user ID is equal to `User_1`:

```
ksql > SELECT * FROM pageviews_original WHERE user_id='User_1' EMIT  
CHANGES;
```

Hit `Ctrl + C` to end the query.

17. Now only records whose page is in the range 60 to 69:

```
ksql > SELECT * FROM pageviews_original  
Ê WHERE page_id LIKE 'Page_6%'  
Ê EMIT CHANGES;
```



The correct filter would be `Page_6_` where we want to only match one character after the 6, but ksqlDB currently only supports the wildcard `%` in the filter, matching zero or more characters. See <https://docs.ksqldb.io/en/latest/developer-guide/ksqldb-reference/select-push-query/#like>

18. We can selectively output information from the stream. Only display the field `user_id` and `page_id`:

```
ksql > SELECT page_id, user_id  
Ê FROM pageviews_original  
Ê EMIT CHANGES  
Ê LIMIT 5;
```

Cleaning Up

19. Exit ksqlDB with `Ctrl +D`.
20. Return to the terminals running ksql-datagen and stop them using `Ctrl +C`.
21. Shut down your Kafka cluster with the `docker-compose down -v` command.

Conclusion

In this lab we have authored our very first ksqlDB queries and run them against two topics `pageview`s and `users` in Kafka. Kafka and all the other components of the ksqlDB platform ran in Docker containers to make the setup very easy and portable.

b. Using the ksqlDB REST API

The idea of this exercise is to show how ksqlDB server can be accessed via its RESTful API by any language that can do `HTTP POST` requests. It is possible to create stream processing applications in any such language by submitting queries to a ksqlDB server cluster. The example in this exercise is a Python application. The Python application is also equipped with Kafka client libraries to produce input data and read output data, but in the real world, there are often dedicated upstream producer and downstream consumer applications, and so the stream processing application would not even need to import Kafka client libraries. ksqlDB enables essentially any language to create real-time stream processing applications via its REST API.

Prerequisites

1. Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)
2. Navigate to the folder `labs/ksql-rest-api`:

```
$ cd ~/confluent-streams/labs/ksql-rest-api
```

3. Run a Kafka cluster and a ksqlDB server using this command:

```
$ docker-compose up -d zookeeper kafka ksqldb-server
Creating network "confluent-streams_kafka-net" with the default
driver
Creating ksqldb-server      ... done
Creating zookeeper          ... done
Creating kafka              ... done
```

■

Wait a couple of minutes until the cluster is ready.

Authoring a Python Client

4. In the folder `labs/ksql-rest-api`, there is a file called `requirements.txt`. It has this content:

```
confluent-kafka==1.4.1
requests
```

These are our external dependencies. The library `confluent-kafka` contains the native Python client for Kafka. The `requests` library we use to `HTTP POST` requests to our ksqldb server

5. Again in the same folder locate and open the file `main.py`. Analyze its content. Discuss it with your peers to make sure you understand what's going on.

In essence the application does the following:

- ! Produce some quotes
- ! Call ksqldb via its REST API to generate a streaming query
- ! Use a Kafka consumer to consume the lowercase quotes produced by the streaming query

6. Create the `quotes` topic:

```
$ kafka-topics \
E  --create \
E  --bootstrap-server kafka:9092 \
E  --replication-factor 1 \
E  --partitions 1 \
E  --topic quotes
```

7. Use `pip3` to install the requirements.

```
$ pip3 install -r requirements.txt
```

#

If you get an error installing the Python packages, run `pip3 install --upgrade pip`. The upgrade will fix the error when you run `pip3 install -r requirements.txt` again.

8. Run the Python client with the following command:

```
$ python3 main.py
```

You should see an output similar to this:

```

>>> Starting Python Kafka Client...
----- Writing quotes to topic 'quotes' -----
*** writing: Kafka enables the Confluent Streaming Platform
*** writing: Confluent offers a Streaming Platform powered by Kafka
*** writing: Kafka Streams are cool
*** writing: Streaming allows for real-time processing of information
*** writing: I love Kafka
----- done writing quotes -----

----- Posting to KSQL Server -----

200, [{"@type": "currentStatus", "statementText": "CREATE STREAM
quotes_orig (line STRING) WITH(KAFKA_TOPIC='quotes',
VALUE_FORMAT='DELIMITED');", "commandId": "stream/`QUOTES_ORIG`/create",
"commandStatus": {"status": "SUCCESS", "message": "Stream
created"}, "commandSequenceNumber": 2, "warnings": []}],

200, [{"@type": "currentStatus", "statementText": "CREATE STREAM
QUOTES_LOWER WITH (KAFKA_TOPIC='QUOTES_LOWER', PARTITIONS=1,
REPLICAS=1) AS SELECT LCASE(QUOTES_ORIG.LINE) KSQL_COL_0\nFROM
QUOTES_ORIG QUOTES_ORIG\nEMIT
CHANGES;", "commandId": "stream/`QUOTES_LOWER`/create", "commandStatus":
{"status": "SUCCESS", "message": "Created query with ID
CSAS_QUOTES_LOWER_0"}, "commandSequenceNumber": 4, "warnings": []}],

----- done posting to KSQL Server -----

>>> Starting Python Kafka Client...
----- Reading from topic 'QUOTES_LOWER' -----
Received message: kafka enables the confluent streaming platform
Received message: confluent offers a streaming platform powered by
kafka
Received message: kafka streams are cool
Received message: streaming allows for real-time processing of
information
Received message: i love kafka
<<< Ending Python Kafka Client...

```

Using the ksqlDB CLI

9. Enter the ksqlDB CLI with this command:

```
$ ksql http://ksql-db-server:8088
```

10. Show all streams:

```
$ ksql > SHOW STREAMS;
```

You should see this:

Stream Name	Kafka Topic	Key Format	Value Format	Windowed
-----	-----	-----	-----	-----
KSQI_PROCESSING_LOG	default_ksql_processing_log	KAFKA	JSON	false
QUOTES_LOWER	QUOTES_LOWER	KAFKA		
DELIMITED	false			
QUOTES_ORIG	quotes	KAFKA		
DELIMITED	false			
-----	-----	-----	-----	-----

11. Use the ksqlDB CLI Print command to list the content of the topic:

```
$ ksql > PRINT 'quotes' FROM BEGINNING;
Key format: ø_(!)_/ø - no data processed
Value format: KAFKA_STRING
rowtime: 2022/04/20 14:15:39.093 Z, key: <null>, value: Kafka enables
the Confluent Streaming Platform, partition: 0
rowtime: 2022/04/20 14:15:39.093 Z, key: <null>, value: Confluent
offers a Streaming Platform powered by Kafka, partition: 0
rowtime: 2022/04/20 14:15:39.093 Z, key: <null>, value: Kafka Streams
are cool, partition: 0
rowtime: 2022/04/20 14:15:39.093 Z, key: <null>, value: Streaming
allows for real-time processing of information, partition: 0
rowtime: 2022/04/20 14:15:39.093 Z, key: <null>, value: I love Kafka,
partition: 0
```

Press **Ctrl +C** to end the query.

12. Set the starting point of your queries to **earliest**:

```
ksql > SET 'auto.offset.reset' = 'earliest';
```

13. Set the column width to **50**:

```
ksql > SET CLI COLUMN-WIDTH 50;
```

14. List the content of both streams:

```
ksql> SELECT * FROM quotes_orig EMIT CHANGES LIMIT 5;
+-----+
| LINE |
+-----+
| Kafka enables the Confluent Streaming Platform |
| Confluent offers a Streaming Platform powered by Kafka |
| Kafka Streams are cool |
| Streaming allows for real-time processing of information |
| I love Kafka |
Limit Reached
Query terminated
```

and

```
ksql> SELECT * FROM quotes_lower EMIT CHANGES LIMIT 3;
+-----+
| KSQL_COL_0 |
+-----+
| kafka enables the confluent streaming platform |
| confluent offers a streaming platform powered by kafka |
| kafka streams are cool |
Limit Reached
Query terminated
ksql>
```

15. Quit the ksqlDB CLI with **Ctrl +D**

Cleanup

16. Shut down your Kafka cluster with the **docker-compose down -v** command.

Conclusion

In this exercise we have created a Kafka client application in Python that uses the ksqldb REST API to access the ksqldb functionality. The Python client executed the following tasks:

- ¥ write entries to an existing topic `quotes`
- ¥ post a query to ksqldb server to create a stream from the topic `quotes`
- ¥ post another query to ksqldb server to create a stream `quotes_lower` containing the quotes from the topic `quotes` all in lower case
- ¥ read from the topic `quotes_lower` and output the messages to the screen

c. Creating connectors with ksqldb

The idea of this exercise is to show how to create connectors using ksqldb server. In the lab, you will create two JDBC source connectors to import a table from PostgresDB to Kafka. One connector will have a simple configuration just to move the data to Kafka; and the other connector will be configured with SMTs (Single Message Transforms) to transform the data while is imported to Kafka.

The architecture of this lab is composed by a Zookeeper cluster, Kafka cluster, ksqldb cluster, Kafka Connect cluster and PostgresDB. ksqldb is configured to communicate to Connect, so you can manage the connectors in Connect using KSQL queries.

Prerequisites

1. Open a terminal window and navigate to the folder `~/confluent-streams`:

```
$ cd ~/confluent-streams
```

2. Run the application with:

```
$ docker-compose up -d zookeeper kafka schema-registry ksqldb-server  
control-center connect postgres
```

3. Create the `postgres-no-key-passengers` topic:


```
$ kafka-topics \
E   --create \
E   --bootstrap-server kafka:9092 \
E   --replication-factor 1 \
E   --partitions 1 \
E   --topic postgres-no-key-passengers
```

4. Create the `postgres-with-key-passengers` topic:

```
$ kafka-topics \
E   --create \
E   --bootstrap-server kafka:9092 \
E   --replication-factor 1 \
E   --partitions 1 \
E   --topic postgres-with-key-passengers
```

Preparing Postgres Database

5. First connect to the Postgres database:

```
$ psql -h postgres -U postgres
psql (11.2)
Type "help" for help.

postgres=#
```

6. At the postgres prompt use a SQL statement to create a new table with the name `passengers`. Run:

```
postgres=# create table passengers("Id" int primary key not null,
"Name" varchar (100), "Email" varchar (255), "Age" integer,
"Travel_to" varchar (255), "Payment" integer, "Travel_date" date);
```

7. Insert some data to the table by running:

```
postgres=# INSERT INTO "passengers" ("Id", "Name", "Email", "Age",
"Travel_to", "Payment", "Travel_date")
VALUES
(1, 'Jack', 'jack12@gmail.com', 20, 'Paris', 79000, '2018-1-1'),
(2, 'Anna', 'anna@gmail.com', 19, 'NewYork', 405000, '2019-10-3'),
(3, 'Wonder', 'wonder2@yahoo.com', 32, 'Sydney', 183000, '2012-8-5'),
(4, 'Stacy', 'stacy78@hotmail.com', 28, 'Maldives', 29000, '2017-6-9'),
(5, 'Stevie', 'stevie@gmail.com', 49, 'Greece', 56700, '2021-12-12'),
(6, 'Harry', 'harry@gmail.com', 22, 'Hogwarts', 670000, '2020-1-17'),
(7, 'Max', 'max@gmail.com', 19, 'Paris', 61000, '2022-4-27');
```

8. Use a SQL select statement to view the contents of the `passengers` table:

```
postgres=# select * from passengers;
```

Id	Name	Email	Age	Travel_to	Payment	Travel_date
1	Jack	jack12@gmail.com	20	Paris	79000	2018-01-01
2	Anna	anna@gmail.com	19	NewYork	405000	2019-10-03
3	Wonder	wonder2@yahoo.com	32	Sydney	183000	2012-08-05
4	Stacy	stacy78@hotmail.com	28	Maldives	29000	2017-06-09
5	Stevie	stevie@gmail.com	49	Greece	56700	2021-12-12
6	Harry	harry@gmail.com	22	Hogwarts	670000	2020-01-17
7	Max	max@gmail.com	19	Paris	61000	2022-04-27

9. Stop the query by typing `:q`.
10. Exit `psql` by pressing `Ctrl +D`.

Install the Kafka Connect JDBC Connector

We use the Kafka Connect JDBC connector in this exercise so we need to install the connector JAR file in Kafka Connect before we can create a JDBC connector.

11. Install the connector:

```
$ docker-compose exec -u root connect confluent-hub install
confluentinc/kafka-connect-jdbc: 10.0.0
The component can be installed in any of the following Confluent
Platform installations:
Ê 1. / (installed rpm/deb package)
Ê 2. / (where this tool is installed)
Choose one of these to continue the installation (1-2):
```

12. Type **1** and press Enter.

13. At the prompt, type **y** and press Enter.

```
Do you want to install this into /usr/share/confluent-hub-components?
(yN)
```

14. At the prompt, type **y** and press Enter.

```
Component's license:
Confluent Community License
https://www.confluent.io/confluent-community-license
I agree to the software license agreement (yN)
```

15. At the prompt, type **y** and press Enter.

```
Downloading component Kafka Connect JDBC 10.0.0, provided by
Confluent, Inc. from Confluent Hub and installing into
/usr/share/confluent-hub-components
Detected Worker's configs:
Ê 1. Standard: /etc/kafka/connect-distributed.properties
Ê 2. Standard: /etc/kafka/connect-standalone.properties
Ê 3. Standard: /etc/schema-registry/connect-avro-
distributed.properties
Ê 4. Standard: /etc/schema-registry/connect-avro-
standalone.properties
Ê 5. Used by Connect process with PID : /etc/kafka-connect/kafka-
connect.properties
Do you want to update all detected configs? (yN)
```

The installation completes.

```
Adding installation directory to plugin path in the following files:  
$ /etc/kafka/connect-distributed.properties  
$ /etc/kafka/connect-standalone.properties  
$ /etc/schema-registry/connect-avro-distributed.properties  
$ /etc/schema-registry/connect-avro-standalone.properties  
$ /etc/kafka-connect/kafka-connect.properties
```

Completed

16. To complete the installation, we need to restart the **connect** container:

```
$ docker-compose restart connect
```

17. Verify that the Connect Worker successfully restarted prior to continuing to the next step:

```
$ docker-compose logs connect | grep -i "INFO .* Finished starting  
connect | [2022-05-13 19:46:25,684] INFO [Worker clientId=connect-1,  
groupId=connect] Finished starting connectors and tasks  
(org.apache.kafka.connect.runtime.distributed.DistributedHerder: 1236)  
connect | [2022-05-13 20:04:11,726] INFO [Worker clientId=connect-1,  
groupId=connect] Finished starting connectors and tasks  
(org.apache.kafka.connect.runtime.distributed.DistributedHerder: 1236)
```



Repeat this command until the Finished starting connectors and tasks message appears.

Create basic connector using ksqlDB

18. Go to Control Center <http://localhost:9021> and click on ksqlDB at the left pane.
19. Click on the ksqlDB application called ksqldb.
20. Run the following KSQL query in the Editor box to create a JDBC source connector:

```
CREATE SOURCE CONNECTOR JDBC_SOURCE_POSTGRES_NO_KEY WITH (
  'connector.class' =
  'io.confluent.connect.jdbc.JdbcSourceConnector',
  'connection.url' = 'jdbc:postgresql://postgres:5432/postgres',
  'connection.user' = 'postgres',
  'table.whitelist' = 'passengers',
  'mode' = 'incrementing',
  'incrementing.column.name' = 'Id',
  'topic.prefix' = 'postgres-no-key-'
);
```



This connector will import each row as a single message in Kafka. By default, this connector puts all values of each row in the Kafka's message value, leaving the message key **null**.

21. Wait for a few seconds until the new connector starts importing the data. Navigate to Topics and select postgres-no-key-passengers topic.
22. In the Overview tab, there is a table with one partition at the bottom of the page. In the column Offset, you should see **Start = 0** and **End = 7** indicating that seven messages have been produced to this topic.
23. In the Schema tab, you can see the AVRO schema automatically generated by the JDBC connector based on the metadata of the Postgres table.

Create connector with SMTs using ksqlDB

24. Navigate again to ksqlDB in Control Center to create a second connector with Single Message Transforms (SMTs).
25. Run the following KSQL query in the Editor box to create a JDBC source connector with SMTs:

```

CREATE SOURCE CONNECTOR JDBC_SOURCE_POSTGRES_WITH_KEY WITH (
  Ê   'connector.class' =
  'io.confluent.connect.jdbc.JdbcSourceConnector' ,
  Ê   'connection.url' = 'jdbc:postgresql://postgres:5432/postgres' ,
  Ê   'connection.user' = 'postgres' ,
  Ê   'table.whitelist' = 'passengers' ,
  Ê   'mode' = 'incrementing' ,
  Ê   'incrementing.column.name' = 'Id' ,
  Ê   'topic.prefix' = 'postgres-with-key-' ,
  Ê   'transforms' =
  'copyFieldToKey,extractKeyFromStruct,removeKeyFromValue' ,
  Ê   'transforms.copyFieldToKey.type' =
  'org.apache.kafka.connect.transforms.ValueToKey' ,
  Ê   'transforms.copyFieldToKey.fields' = 'Id' ,
  Ê   'transforms.extractKeyFromStruct.type' =
  'org.apache.kafka.connect.transforms.ExtractField$Key' ,
  Ê   'transforms.extractKeyFromStruct.field' = 'Id' ,
  Ê   'transforms.removeKeyFromValue.type' =
  'org.apache.kafka.connect.transforms.ReplaceField$Value' ,
  Ê   'transforms.removeKeyFromValue.blacklist' = 'Id' ,
  Ê   'key.converter' =
  'org.apache.kafka.connect.converters.IntegerConverter'
);

```

Explanation about the new code:

!

1. There are three transforms to set the key to the value of the **Id** field.

They run in the order defined by **transforms**:

- " **copyFieldToKey** sets the key to a struct containing the **Id** field from the value.
- " **extractKeyFromStruct** sets the key to just the **Id** field of the struct set by the previous step.
- " **removeKeyFromValue** removes the **Id** field from the message value, as it's now stored in the message key.

2. Last line - the key is an integer so we override the default serialization (**StringConverter**) and instead use the **IntegerConverter** for the key field.

26. Wait for a few seconds until the new connector starts importing the data. Navigate to Topics and select postgres-with-key-passengers topic.

27. In the Overview tab, you should see `Start = 0` and `End = 7` in the column Offset, indicating that seven messages have been produced to this topic.
28. In the Schema tab, you can see the AVRO schema automatically generated by the JDBC connector. Note that the field `Id` is not present in the value schema, since it has been moved to the message key by the SMTs.

Read topics using ksqlDB

In this part, we will read the data imported to `postgres-no-key-passengers` and `postgres-with-key-passengers` topics.

29. Navigate to ksqlDB in Control Center.
30. Let's now inspect the data on the Kafka topic `postgres-no-key-passengers`. ksqlDB's `PRINT` command will show the contents of a topic. Run this query:

```
PRINT 'postgres-no-key-passengers' FROM BEGINNING;
```



Note that `key: <null>` and the field `Id` is within the value.

31. Click `Stop` to finish the query.
32. Now, read the messages in `postgres-with-key-passengers` topic by running this query:

```
PRINT 'postgres-with-key-passengers' FROM BEGINNING;
```



Check the results. In this case the `key` is the field `Id` value.

33. Click `Stop` to finish the query.

Optional: Extra Content

You can use the `kafka-avro-console-consumer` to read AVRO data using the Terminal. This tool uses the Avro converter with the Schema Registry in order to properly read the Avro data schema and write the messages to standard output (console) in JSON format.

- a. Open a Terminal window.
- b. Run this command to read the messages from `postgres-no-key-passengers` topic:

```
$ kafka-avro-consol e-consumer \  
Ê --bootstrap-server kafka: 9092 \  
Ê --topi c postgres-no-key-passengers \  
Ê --from-begi nni ng \  
Ê --property schema. regi stry. url =http: //schema-regi stry: 8081 \  
Ê --property print. key=true \  
Ê --property key. separator=" | "  
null |  
{ "Id": 1, "Name": { "stri ng": "Jack" }, "Email": { "stri ng": "j ack12@gmai l . com"  
, "Age": { "i nt": 20 }, "Travel _to": { "stri ng": "Pari s" }, "Payment": { "i nt": 79  
000 }, "Travel _date": { "i nt": 17532 } }  
null |  
{ "Id": 2, "Name": { "stri ng": "Anna" }, "Email": { "stri ng": "anna@gmai l . com" },  
"Age": { "i nt": 19 }, "Travel _to": { "stri ng": "NewYork" }, "Payment": { "i nt": 40  
5000 }, "Travel _date": { "i nt": 18172 } }  
null |  
{ "Id": 3, "Name": { "stri ng": "Wonder" }, "Email": { "stri ng": "wonder2@yahoo. c  
om" }, "Age": { "i nt": 32 }, "Travel _to": { "stri ng": "Sydney" }, "Payment": { "i nt  
": 183000 }, "Travel _date": { "i nt": 15557 } }  
...
```

- c. Press `Ctrl +C` to stop the `kafka-avro-consol e-consumer`.
- d. Run this command to read the messages from `postgres-no-key-passengers` topic:


```
$ kafka-avro-console-consumer \
  --bootstrap-server kafka:9092 \
  --topic postgres-with-key-passengers \
  --from-beginning \
  --property schema.registry.url=http://schema-registry:8081 \
  --property print.key=true \
  --property key.separator=" | " \
  --key-deserializer
org.apache.kafka.common.serialization.IntegerDeserializer
1 |
{"Name": {"string": "Jack"}, "Email": {"string": "jack12@gmail.com"}, "Age": {"int": 20}, "Travel_to": {"string": "Paris"}, "Payment": {"int": 79000}, "Travel_date": {"int": 17532}}
2 |
{"Name": {"string": "Anna"}, "Email": {"string": "anna@gmail.com"}, "Age": {"int": 19}, "Travel_to": {"string": "NewYork"}, "Payment": {"int": 405000}, "Travel_date": {"int": 18172}}
3 |
{"Name": {"string": "Wonder"}, "Email": {"string": "wonder2@yahoo.com"}, "Age": {"int": 32}, "Travel_to": {"string": "Sydney"}, "Payment": {"int": 183000}, "Travel_date": {"int": 15557}}
...
```



Note we are overriding the default `--key-deserializer` with `IntegerDeserializer`, since the key was serialized as `Integer` by the JDBC connector.
Default deserializer is `StringDeserializer`.

e. Press **Ctrl + C** to stop the `kafka-avro-console-consumer`.

Manage connectors using ksqlDB

34. Check all connectors that you have created. Run this query in the Editor box in ksqlDB:

```
SHOW CONNECTORS;
```

In the result, you should see two connectors running:

```
! JDBC_SOURCE_POSTGRES_NO_KEY
```

```
! JDBC_SOURCE_POSTGRES_WITH_KEY
```

35. Get extra information about a connector using the function `DESCRIBE CONNECTOR`. Run

this query:

```
DESCRIBE CONNECTOR JDBC_SOURCE_POSTGRES_WITH_KEY;
```

36. Use the function `DROP CONNECTOR` to delete a connector from the Connect cluster. The topics associated with this connector are not deleted by this command. Run this query to delete `JDBC_SOURCE_POSTGRES_NO_KEY` connector:

```
DROP CONNECTOR JDBC_SOURCE_POSTGRES_NO_KEY;
```

37. Check the connector was deleted by running again:

```
SHOW CONNECTORS;
```

Cleanup

38. Shut down all Docker containers with the `docker-compose down -v` command.

Conclusion

In this lab, you have learnt how to create connectors using ksqlDB, how to read data from Kafka topics using ksqlDB and how to manage connectors using ksqlDB. Additionally, you have learnt how to install a connector in a Kafka Connect cluster and how to use the tool `kafka-avro-console-consumer` to read AVRO data from Kafka using your console.

STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 04 Using ksqlDB

a. Using ksqlDB

In this exercise, we will explore how to apply ksqlDB's mapping and filtering capabilities to an application that processes real-time temperature data. MQTT and Internet of Things data are perfect for real-time processing with Kafka Streams and ksqlDB. ksqlDB is an especially good choice for many of these applications because of its simplicity.

Prerequisites

Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)

Preparing the Kafka Cluster

1. Open a terminal window and navigate to the folder `~/confluent-streams`

```
Ê $ cd ~/confluent-streams
```

2. Run a Kafka cluster and a ksqlDB server using this command:

```
$ docker-compose up -d zookeeper kafka ksql db-server
Creating network "confluent-streams_kafka-net" with the default driver
Creating ksql db-server      ... done
Creating zookeeper          ... done
Creating kafka               ... done
```



Wait a couple of minutes until the cluster is ready.

3. Check the status with:

```
$ docker-compose ps
```

Creating Data

4. Create the **stations** topic:

```
$ kafka-topics \
Ê --create \
Ê --bootstrap-server kafka:9092 \
Ê --replication-factor 1 \
Ê --partitions 1 \
Ê --topic stations \
Ê --config cleanup.policy=compact
```



We do this compaction only for illustration. Normally the temperature reading stations wouldn't change too frequently to warrant compaction.

5. Create a list of stations. Here we use the kafka-console-producer command line tool to send records in to the Kafka cluster.

```
$ cat << EOF | kafka-console-producer \
Ê      --bootstrap-server kafka:9092 \
Ê      --property "parse.key=true" \
Ê      --property "key.separator=: " \
Ê      --topic stations
1: Mombasa, Kenya
2: Nairobi , Kenya
3: Mogadishu, Somalia
4: Dar es Salaam, Tanzania
5: Pretoria, South Africa
6: Cape Town, South Africa
7: Bloemfontein, South Africa
8: Diani , Kenya
9: Embu, Kenya
10: Johannesburg, South Africa
EOF

>>>>>>>>>>
```

6. Double check that the list of stations has been created:

```
$ kafka-console-consumer \
  --bootstrap-server kafka:9092 \
  --from-beginning \
  --max-messages 7 \
  --topic stations \
  --property print.key=true \
  --property key.separator=": "
1: Mombasa, Kenya
2: Nairobi, Kenya
3: Mogadishu, Somalia
4: Dar es Salaam, Tanzania
5: Pretoria, South Africa
6: Cape Town, South Africa
7: Bloemfontein, South Africa
8: Diani, Kenya
9: Embu, Kenya
10: Johannesburg, South Africa
Processed a total of 10 messages
```

7. Create the **temperatures** topic:

```
$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic temperatures
```

8. Create a list of temperature readings, again using the kafka-console-producer command line tool.

[illegible]

Running the ksqlDB CLI

9. Open a new terminal window and run the `ksqlDB` CLI using this command:

```
$ ksql http://ksql db-server: 8088
```

10. Define that streams should be read from beginning:

```
ksql> SET 'auto.offset.reset' = 'earliest';  
Successfully changed local property 'auto.offset.reset' from 'null'  
to 'earliest'
```

Mapping and Filtering

11. Create a table from the topic `stations`:

```
ksql > CREATE TABLE weather_stations(  
  id VARCHAR PRIMARY KEY,  
  name VARCHAR,  
  country VARCHAR  
)  
WITH(kafka_topic='stations', value_format='DELIMITED');
```

Message

Table created

12. Run a simple query against this new table:

```
ksql > SELECT * FROM weather_stations EMIT CHANGES LIMIT 5;
```

you should see something like this:

```
+-----+-----+-----+  
| ID    | NAME          | COUNTRY    |  
+-----+-----+-----+  
| 1     | Mombasa       | Kenya    |  
| 2     | Nairobi       | Kenya    |  
| 3     | Mogadishu     | Somalia    |  
| 4     | Dar es Salaam | Tanzania   |  
| 5     | Pretoria      | South Africa |  
Limit Reached  
Query terminated
```

You can also peek into the `stations` topic using:

!

```
ksql > PRINT 'stations' FROM BEGINNING;
```

Press `Ctrl +C` to stop the above query

13. Now let's only output stations in Kenya:

```
ksql> SELECT * FROM weather_stations
Ê     WHERE country='Kenya'
Ê     EMI T CHANGES;
```

ID	NAME	COUNTRY
1	Mombasa	Kenya
2	Nai robi	Kenya
8	Di ani	Kenya
9	Embu	Kenya



Press **Ctrl +C** to end the query.

14. To show only stations whose name starts with M and output the country in all caps use:

```
ksql> SELECT id, name, UCASE(country) AS country
Ê     FROM weather_stations
Ê     WHERE name LIKE 'M%'
Ê     EMI T CHANGES;
```

ID	NAME	COUNTRY
1	Mombasa	Kenya
2	Mogadi shu	Somal i a

Press **Ctrl +C** to end the query.

15. Create a stream from the topic **temperatures**:

```
ksql> CREATE STREAM mytemperatures(
Ê     id INTEGER,
Ê     station_id VARCHAR,
Ê     temp DOUBLE
Ê )
Ê WITH(kafka_topi c=' temperatures', val ue_format=' DELI MI TED' );
```

16. The temperatures are in degree Celsius. To output them in degree Fahrenheit use this:


```
ksql> SELECT id, station_id,
           temp AS temp_in_C,
           temp*9/5+32 as temp_in_F
        FROM mytemperatures
        EMIT CHANGES
        LIMIT 10;
```

ID	STATION_ID	TEMP_IN_C	TEMP_IN_F
6	2	18.5	65.3
5	1	23.0	73.4
20	5	21.3	70.34
15	4	35.5	95.9
17	4	34.5	94.1
9	2	18.0	64.4
24	7	17.0	62.6
13	3	33.0	91.4
19	5	21.0	69.8
25	7	18.0	64.4

```
Limit Reached
Query terminated
```

Cleanup

- Exit the ksqlDB CLI by pressing **Ctrl +D**.
- Shut down your Kafka cluster with the **docker-compose down -v** command.

Conclusion

In this exercise, we have explored the various capabilities that ksqlDB offers us in a easy and convenient way. We have learned that the syntax of ksqlDB SQL strongly resembles that of ANSI SQL. We have explored stateless functions such as mapping and filtering.

STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 06 Windowing & Aggregations

In this exercise you will create a Kafka Streams application and a ksqlDB application to sessionize click data from a website. You will be organizing the user's click behaviour data collected from a website using sessions window.

You can choose to work on either Kafka Streams or ksqlDB lab or both.

a. Windowing & Aggregations

Prerequisites

Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)

Writing the data to Kafka

1. Use the command in the table below to navigate to the project folder for your language:

```
cd ~/confluent-streams/labs/windowing
```

2. If your Kafka cluster is not already running, start it with:

```
$ docker-compose up -d zookeeper kafka schema-registry ksqldb-server control-center
```

3. Create one input topic called `clicks-topic` and two output topics called `window-streams` and `window-ksql` in Kafka:

```

$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic clicks-topic

$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic window-streams

$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic window-ksql

```

4. Start producing data to Kafka using the Java producer `clicks-producer`:

a. Navigate to the producer folder:

```
$ cd ./protobuf-java-producer
```

b. Use `gradle` to generate the PROTOBUF class:

```
$ ./gradlew build
```

c. Run this command to start the Java producer:

```
$ ./gradlew run
```

5. We have the clicks data in `clicks-topic`. Now, we are going to use Kafka Streams and ksqlDB to do the same operation (use session windows to count the number of clicks for each IP address).

Using Session Windows in Kafka Streams

6. Open a new Terminal window and run the following command to open count-streams-app in Visual Studio Code:

```
$ code ~/confluent-streams/labs/windowing/count-streams-app
```

7. In VS Code, navigate to `src/main/java/io/confluent/training/app/` and open the Java file `StreamsApp.java`
8. Locate in the code the `TO-DO` lines and `{{ WRITE-MISSING-CODE }}` markers. Try to write yourself the missing code by checking the documentation:

! TO-DO 1 - [documentation](#): create a `KStream` from the "clicks-topic" topic and configure the Key-Serde and Value-Serde that can read the String key, and Clicks value.

Solution

```
final KStream<String, ClicksProtos.Clicks> clicks =  
builder.stream("clicks-topic", Consumed.with(Serdes.String(),  
clicksSerde));
```

! TO-DO 2 - [documentation](#): group by key the `KStream` "clicks".

Solution

```
final KGroupedStream<String, ClicksProtos.Clicks> clicksGrouped  
= clicks.groupByKey();
```

! TO-DO 3 - [documentation](#): apply a Session Window of 5 minutes with a Grace period of 30 seconds to the `KGroupedStream` "clicksGrouped" and apply a count to get the number of clicks per IP per Session Window.

Solution

```
final KTable<Windowed<String>, Long> clicksCount =  
clicksGrouped  
    .windowedBy(SessionWindows.with(Duration.ofMinutes(5)).g  
        race(Duration.ofSeconds(30)))  
    .count();
```

! TO-DO 4 - [documentation](#): convert the `KTable` "clicksCount" into a `KStream`.

Solution

```
final KStream<Windowed<String>, Long> clicksCountStream =  
clicksCount.toStream();
```

- ! TO-DO 5 - [documentation](#): produce the data of the KStream "clicksCountStreamModified" to the topic "window-streams" selecting the appropriate Serdes for key and value.

Solution

```
clicksCountStreamModified.to("window-streams",  
Produced.with(Serdes.String(), clicksCountSerde));
```

9. After completing all the TO-DOs, start the Kafka Streams application by clicking in the top menu: Run ! Start Debugging.
10. Let the application run for a few seconds and check the results in Confluent Control Center <http://localhost:9021>:
 - ! Go to Topics and select window-streams topic
 - ! Click on the tab Messages
 - ! In the box at the top where says **offset** with a magnifier, type **0**. A dropdown list will appear, then select **0/Partition: 0**
 - ! Wait for a few seconds and you will see the output from the Kafka Streams app
11. Check the results and try to answer the following questions:
 - ! How many different IP addresses are there?
 - ! How many different sessions are there?
 - ! Which IP has the highest number of clicks?



It might be easier to answer those questions when "Table View" is selected for the messages in top right corner in Control Center, as compared to the "Cards View".

Using Session Windows in ksqlDB

In this section we are going to apply the same transformations but using ksqlDB, so you can easily compare the differences and similarities between both approaches.

12. Go to Control Center <http://localhost:9021> and click on ksqlDB at the left pane.
13. Click on the ksqlDB application called ksqldb. This is your ksqlDB cluster formed by one server which is the ksqldb-server Docker container.
14. Set the following property from dropdown to ensure that you're reading from the beginning:

```
auto.offset.reset = Earliest
```

15. Now you will create and populate a new stream `clicks` with the data from `clicks-topic`. An important characteristic of this data is the timestamp because this is what drives the session window. ksqlDB can use either the Kafka message timestamp, or a field from the message value as the timestamp. In this example we'll use the latter—the event time as stored in the `timestamp` field of the message value. Copy and paste the following code in the Editor box, then click on Run query:

```
CREATE STREAM clicks
  WITH (KAFKA_TOPIC='clicks-topic',
  TIMESTAMP='timestamp',
  VALUE_FORMAT='PROTOBUF');
```

16. Run the following non-persistent query to check that the pseudo column `ROWTIME` (system column) contains the same value in milliseconds as the message column `timestamp`:

```
SELECT ROWTIME, timestamp FROM clicks EMIT CHANGES LIMIT 5;
```

17. If you want to see the timestamp values in a more readable format, use the function `TIMESTAMP ToString()`. Run this query:

```
SELECT TIMESTAMPToString(ROWTIME, 'yyyy-MM-dd HH:mm:ss', 'UTC') AS
ROWTIME_STR, TIMESTAMPToString(timestamp, 'yyyy-MM-dd HH:mm:ss',
'UTC') AS TIMESTAMP_STR FROM clicks EMIT CHANGES LIMIT 5;
```

18. Before running the `COUNT()` aggregation, configure ksqlDB to buffer the aggregates as it builds them. This makes the query feel like it responds more slowly, but it means that you get just one row per window. This makes it simpler to understand the results:

- ! Click on +Add another field

- ! Type: `ksql.streams.cache.max.bytes.buffering = 2000000`

19. Using the [documentation](#), write a non-persistent query with Session window to count how many clicks were made in each user session (based on IP address). Set the Session window with a gap of 5 minutes and a grace period of 30 seconds. The output should contain these two columns:

- ! IP

- ! CLICK_COUNT

Solution

```
SELECT i p,  
       COUNT(*) AS CL I CK_COUNT  
FROM c l i c k s  
   WINDOW SESSION (5 MI NUTES, GRACE PERIOD 30 SECONDS)  
GROUP BY i p  
EMIT CHANGES;
```

20. From the previous non-persistent query, add three more columns to the result to include the start window timestamp, the end window timestamp and the window length in milliseconds. Use the function `TI MESTAMPTOSTRI NG()` to print nicely the start/end window timestamps. Output columns:

- ! IP

- ! CLICK_COUNT

- ! SESSION_START_TS

- ! SESSION_END_TS

- ! SESSION_LENGTH_MS



When you apply a **WINDOW** clause, ksqlDB adds two additional system columns to the data, which provide the window bounds: **WINDOWSTART** and **WINDOWEND**.

Solution

```
SELECT ip,
       COUNT(*) AS CLICK_COUNT,
       TIMESTAMPTOSTRING(WINDOWSTART, ' yyyy-MM-dd HH: mm: ss' ,
       ' UTC' ) AS SESSION_START_TS,
       TIMESTAMPTOSTRING(WINDOWEND, ' yyyy-MM-dd HH: mm: ss' , ' UTC' )
AS SESSION_END_TS,
       WINDOWEND - WINDOWSTART AS SESSION_LENGTH_MS
FROM clicks
       WINDOW SESSION (5 MINUTES, GRACE PERIOD 30 SECONDS)
GROUP BY ip
EMIT CHANGES;
```

21. Finally, write a persistent query based on the previous query creating a new table called **IP_SESSIONS** using as backing topic **window-ksql**.

Solution

```
CREATE TABLE IP_SESSIONS
       WITH (KAFKA_TOPIC=' window-ksql ' ) AS
SELECT ip,
       COUNT(*) AS CLICK_COUNT,
       TIMESTAMPTOSTRING(WINDOWSTART, ' yyyy-MM-dd HH: mm: ss' , ' UTC' )
AS SESSION_START_TS,
       TIMESTAMPTOSTRING(WINDOWEND, ' yyyy-MM-dd HH: mm: ss' , ' UTC' )
AS SESSION_END_TS,
       WINDOWEND - WINDOWSTART AS SESSION_LENGTH_MS
FROM clicks
       WINDOW SESSION (5 MINUTES, GRACE PERIOD 30 SECONDS)
GROUP BY ip;
```

22. Check the results in Confluent Control Center <http://localhost:9021>:

- ! Go to ksqlDB at the left pane, click on ksqldb application and click on the tab Persistent queries to analyze the query
- ! Click on the tab Flow to visualize the data flow in your ksqlDB app

- ! Go to Topics at the left pane and select window-ksql topic
- ! Click on the tab Messages
- ! In the box at the top where says `offset` with a magnifier, type `0`. A dropdown list will appear, then select `0/Partition: 0`
- ! Wait for a few seconds and you will see the output from your ksqlDB app

Cleanup

1. Terminate your Kafka Streams application with Run ! Stop Debugging.
2. Shut down your Kafka cluster with the `docker-compose down -v` command.

Conclusion

In this exercise you created a Kafka Streams application and a ksqlDB application to sessionize click data from a website. A given user might visit a website multiple times a day, but in distinct visits; so using Session window you could automatically organize the data in sessions based on a period of inactivity.

You also learned how to use the ksqlDB system columns (`ROWTIME`, `WINDOWSTART`, `WINDOWEND`) and to implement time extractors in Kafka Streams and ksqlDB queries.

STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 07 Joins

a. Joining Two Streams

In event-driven architecture, it is important to think about what event will trigger an output. Different kinds of joins will trigger outputs under different conditions. The purpose of this exercise is to experiment with the behavior of various stream-stream joins.

The streaming application itself performs a stream-stream join. It takes a string value from a "left stream" and a string value from a "right stream" and concatenates them together inside of brackets. The output is produced to a new stream. Remember that all stream-stream joins must be windowed since streams are unbounded. This application will use a tumbling window of 5 minutes.

Prerequisites

1. Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)
2. If your Kafka cluster is not already running, start it with:

```
$ cd ~/confluent-streams/labs/joining-streams
$ docker-compose up -d zookeeper kafka control-center ksqldb-server
schema-registry
```

3. Create two input topics called `left-topic` and `right-topic` and an output topic called `joined-topic` in Kafka:

```

$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic left-topic

$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic right-topic

$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic joined-topic

```



Remember that joins require the input topics to have the same number of partitions so that all input records with the same key, from both sides of the join, are delivered to the same stream task during processing. (called co-partitioning).

Create the Streaming App

4. Open another terminal and navigate to the `joining-streams` folder, and launch VS Code:

```

$ cd ~/confluent-streams/abs/joining-streams
$ code .

```

5. Open the file `build.gradle` and analyze its content.
6. Locate the file `JoinSample.java` in the folder `src/main/java/streams` and open it. Familiarize yourself with the code. It basically does the configuration of the Kafka Streams app but the interesting code is missing. If you would like to challenge yourself, take this time to create the streaming application logic yourself.
7. Now we define the Topology for our application. Add the following code snippet to the

`getTopology()` function after the TODO comment:

```
KStream<String, String> leftStream = builder.stream("left-topic",
    Consumed.with(stringSerde, stringSerde));
KStream<String, String> rightStream = builder.stream("right-topic",
    Consumed.with(stringSerde, stringSerde));
leftStream
    .join(rightStream,
        (leftValue, rightValue) -> "[" + leftValue + ", " +
            rightValue + "]",
        JoinWindows.of(Duration.ofMinutes(5)),
        StreamJoined.with(stringSerde, stringSerde, stringSerde)
    )
    .to("joined-topic", Produced.with(stringSerde, stringSerde));

Topology topology = builder.build();
return topology;
```

What does the above code do? Discuss with your peers.

- Note the file `log4j.properties` in the folder `src/main/resources`, which is used to configure logging for our streams application.
- Use Run ! Start Debugging to run your code. Let your Kafka Streams app remain running in the VS Code debugger.

You should get this output:

```
*** Starting Join Sample Application ***
...
```



Ignore the WARNINGS.

Creating input Data

- Open 3 terminal windows and arrange them side by side so you can see all three of them at the same time.

We will be using the tool `kafkacat` to generate data and monitor the output:

- In the first terminal window start the tool `kafkacat` as a producer for the `left-topic`

topic:

```
kafkacat \  
-b kafka:9092 \  
-t left-topic \  
-P -K: -Z
```

12. In the second terminal window start `kafkacat` as a producer for the `right-topic` topic:

```
kafkacat \  
-b kafka:9092 \  
-t right-topic \  
-P -K: -Z
```

13. In the third terminal window run an instance of `kafkacat` as a consumer of the `joined-topic` topic:

```
kafkacat \  
-b kafka:9092 \  
-t joined-topic \  
-C -K\\t
```

14. In window 1 enter `FL:` (a record key of `FL` for Florida a `<NULL>` value for the record value) and observe the output in window 3. Hint: Nothing should happen. Why?
15. In window 2 also enter the value `FL:` and observe the output in window 3. Hint: Nothing should happen. Why?
16. Now in window 1 enter the value `FL: Orlando` and observe the output in window 3. Hint: Nothing should happen. Why?
17. In window 2 enter the value `FL: Tampa` and observe the output in window 3. You should see:

```
FL [Orlando, Tampa]
```

18. Back in window 1 enter `FL:` and observe the output in window 3. What do you see? Why?
19. Still in window 1 enter `FL: Miami` and observe the output in window 3. You should see:

```
FL [Miami, Tampa]
```

20. Continue with window 2 and value **FL: Naples**. What do you see this time?
21. Continue to enter more values with the same key **FL**. Here are other cities in Florida to experiment with:
- ! Jacksonville
 - ! Alachua
 - ! Pensacola
 - ! Destin
 - ! Fort Meyers
22. What happens if you use a different key, say **NY**? Why? Discuss the results with your peers.
23. What happens when an event falls outside of the tumbling window?

Optional: Extra Content for ksqlDB INNER Join

24. Go to Control Center <http://localhost:9021> and click on ksqlDB at the left pane to open your **ksql db** application.
25. In the Editor box, paste the following queries to create the ksqlDB streams with the data from the **left-topic** and **right-topic**, and then click Run query:

```

CREATE STREAM left_stream_kafka (
  left_key STRING KEY,
  left_value STRING)
WITH (
  KAFKA_TOPIC='left-topic',
  VALUE_FORMAT='kafka');

CREATE STREAM left_stream
WITH(
  VALUE_FORMAT='avro')
AS SELECT * FROM left_stream_kafka;

CREATE STREAM right_stream_kafka (
  right_key STRING KEY,
  right_value STRING)
WITH (
  KAFKA_TOPIC='right-topic',
  VALUE_FORMAT='kafka');

CREATE STREAM right_stream
WITH(
  VALUE_FORMAT='avro')
AS SELECT * FROM right_stream_kafka;

```

!

We need to duplicate the streams because **KAFKA** format does not yet support **JOIN**. The **KAFKA** format is primarily intended for use as a key format. It can be used as a value format, but can not be used in any operation that requires a repartition or changelog topic. Removing this limitation requires enhancements to the core of KSQL. This will come in a future release. Until then, avoid using the **KAFKA** format for values.

26. Using the documentation about [JOIN](#) and [WITHIN](#), try to write the KSQL query to perform the same join operation you have done in the previous section using a Kafka Streams application.

Solution

```
SELECT
  left_key AS key,
  '[' + left_value + ', ' + right_value + ']' AS joined_value
FROM left_stream
JOIN right_stream
  WITHIN 5 MINUTES
  ON left_key = right_key
EMI T CHANGES;
```

Left Join

27. Modify the `getTopology()` function and replace the `join` function with a `leftJoin` function instead.
28. Recompile and run the application.
29. This time, use a car brand for the key and different car models for values to play with input data as you have done for the (inner) join example.

What is different? Discuss with your peers if needed.

Optional: Extra Content for ksqlDB LEFT Join

30. Go to Control Center <http://localhost:9021> and click on ksqlDB at the left pane to open your `ksql db` application.
31. Try to write the KSQL query to perform the left join operation.

Solution

```
SELECT
  left_key AS key,
  '[' + left_value + ', ' + right_value + ']' AS joined_value
FROM left_stream
LEFT JOIN right_stream
  WITHIN 5 MINUTES
  ON left_key = right_key
EMI T CHANGES;
```

Outer Join

32. Modify the `getTopology()` function and replace the current join function with an `outerJoin` function instead.
33. Recompile and run the application.
34. This time, use your home country for key and different cities for values to play with input data as you have done for the (inner) join example and observe the output.

What is different? Discuss with your peers if needed.

Optional: Extra Content for ksqlDB OUTER Join

35. Go to Control Center <http://localhost:9021> and click on ksqlDB at the left pane to open your `ksql db` application.
36. Try to write the KSQL query to perform the outer join operation.



TIP: Check the function `IFNULL`, you may need to use it.

Solution

```
SELECT
  IFNULL(left_key, right_key) AS key,
  '[' + left_value + ', ' + right_value + ']' AS joined_value
FROM left_stream
FULL JOIN right_stream
  WITHIN 5 MINUTES
  ON left_key = right_key
EMI T CHANGES;
```

Optional: Stream - Table Joins

37. Perform a similar experiment with a stream - table left join (the most common join in most streaming applications). Make sure to experiment with sending null keys and values. How are the results different from the stream - stream left join?

Cleanup

38. Terminate the first 2 instances of `kafkacat` (the producer instances) by pressing `Ctrl +D`.
39. Terminate the third instance of `kafkacat` (the consumer instance) by pressing `Ctrl +C`.
40. Terminate your Kafka Streams application with Run ! Stop Debugging.
41. Shut down your Kafka cluster with the `docker-compose down -v` command.

Conclusion

In this exercise you created a Kafka Streams application and KSQL queries that join two streams with inner, left, and outer joins. You then created data for the left and the right input stream and observed the generated output. You used the command line tool `kafkacat` to generate input data and observe output data. Consider summarizing your observations and comparing them to the information found here:

<https://docs.confluent.io/current/streams/developer-guide/dsl-api.html#kstream-kstream-join>. Especially focus on the subsection called "Semantics of stream-stream joins" with an illustrative table of the output records that are produced from a join as events flow into the left and right streams.

STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 08 Custom Processing

a. Using the Processor API

The purpose of this exercise is to create an application with the lower-level Processor API. This may be required in applications that require a greater level of control over state store management and more sophisticated application logic than the Kafka Streams DSL can provide.

This exercise will give you experience creating a Kafka Streams application using the DSL and creating a new node in the topology using the Processor API through the `transform()` method.

This application uses a simple source ! word count processor ! sink processing topology. The source node takes in records from an input topic whose values are sentence strings. The word count processor uses each record's value to update its internal state store for word counts (key=word, value=count) and sends that state to the sink processor every second. The sink processor produces the resulting records to an output topic.

Prerequisites

1. Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)
2. If your Kafka cluster is not running already, start it with:

```
$ cd ~/confluent-streams/labs/processor-api  
$ docker-compose up -d zookeeper kafka
```

3. Create two topics called `lines-topic`, and `word-count-topic` in Kafka:

```
$ kafka-topics \
E --create \
E --bootstrap-server kafka:9092 \
E --replication-factor 1 \
E --partitions 1 \
E --topic lines-topic

$ kafka-topics \
E --create \
E --bootstrap-server kafka:9092 \
E --replication-factor 1 \
E --partitions 1 \
E --topic word-count-topic
```

Create the Streaming App

4. Open another terminal and navigate to the `processor-api` folder, and launch VS Code:

```
$ cd ~/confluent-streams/labs/processor-api
$ code .
```

5. Open the file `build.gradle` in the `processor-api` folder and analyze its content. It should be quite familiar by now.
6. Open the file `WordCountTransformer.java` in subfolder `src/main/java/streams`.

Familiarize yourself with the code. We are creating an instance of type `Transformer` and overriding its methods. Thoroughly document what the `init` method does. Pay particular note to the call to the `context.schedule()` method, with its punctuation and the call to `context.forward()`. As always, you can check the corresponding code in the `solutions` folder for feedback.

7. You can challenge yourself to write the `transform` method, or continue with the next step.
8. Make it so the `transform` method of the class has these contents:

```
Long oldValue = this.kvStore.get(word);
if (oldValue == null) {
E this.kvStore.put(word, 1L);
} else {
E this.kvStore.put(word, oldValue + 1L);
}
```

This code gets the correct entry from the keystore and updates it or it creates a new entry in the keystore with a key of the incoming word and a count of 1. It leaves it to the context scheduler to forward the entries in the key value store, producing them to Kafka for durability.

9. Open the file `CustomTransformerApp.java` in the same folder, and familiarize yourself with the code. This code is entirely DSL - It defines the topology, creates the configuration, sets up the shutdown hook and starts the Kafka Streams app.

In the topology, it creates a stream from the input topic, uses the `flatMapValues()` method to break up the lines of input text into individual words, writes a rekeyed stream out to a repartition topic and reads it back in. Then it calls the `transform()` method we created in the `WordCountTransformer` code. And finally, it directs those results to the output topic.

10. Notice the file `log4j.properties` in the folder `src/main/resources` used to configure logging for the application.
11. Use Run ! Start Debugging to run your code. Let your Kafka Streams app remain running in the VS Code debugger.

You should get this output:

```
*** Starting Custom Transformer App Application ***
```

Creating input Data

12. Open 2 terminal windows and arrange them side by side so you can see the two at the same time.
13. In the first terminal window start the tool `kafkacat` as a producer for the `lines-topic` topic:

```
kafkacat \
-b kafka:9092 \
-t lines-topic \
-P -K ;
```

14. In the second terminal window run an instance of `kafka-console-consumer` as a consumer of the `word-count-topic` topic, printing the String key and the Long value:

```
kafka-console-consumer --bootstrap-server kafka:9092 \
  --topic word-count-topic --from-beginning \
  --property print.key=true \
  --value-deserializer
org.apache.kafka.common.serialization.LongDeserializer
```

15. In window 1 enter each of these strings one at a time:

```
kafka: Kafka powers the Confluent streaming platform
kafka: A streaming application uses Kafka
kafka: Everyone Loves Kafka
kafka: Many contributors to Kafka work for Confluent
```

After each line observe the output in window 2.

16. Discuss the result with your peers.

Cleanup

17. Terminate the `kafkacat` instance by pressing `Ctrl +D`.
18. Terminate the `kafka-console-consumer` instance by pressing `Ctrl +C`.
19. Terminate your Kafka Streams application with Run ! Stop Debugging.
20. Shut down your Kafka cluster with the `docker-compose down -v` command.

Conclusion

In this sample we have demonstrated the use of a custom transform written using the Processor API to count words. This is included into a Stream processing application written entirely in the Streams DSL that takes in sentences from an input topic, process them and writes the resulting word and count pairs to an output topic.

STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 09 Testing, Monitoring and Troubleshooting

This lab contains 5 exercises:

¥ Testing:

- ! Building Unit Tests
- ! Integration Tests using Embedded Kafka

¥ Monitoring:

- ! Getting Metrics from a Kafka Streams Application
- ! Using JConsole to monitor a Streams App
- ! Monitoring a Kafka Streams App in Confluent Control Center

a. Building Unit Tests

It is essential for every application to have full test coverage of each of its components. The purpose of this exercise is to build unit tests for an existing Kafka Streams application and test the application with Gradle.

Prerequisites

1. Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)
2. Notice that we do not need to run our Kafka cluster - the unit and integration testing do not require it.
3. Navigate to the folder `~/confluent-streams/labs/testing/simple-test`, and launch VS Code:

```
$ cd ~/confluent-streams/labs/testing/simple-test
$ code .
```

4. Open the file `build.gradle` inside this folder and analyze its content. Notice the `junitt`, as well as the `kafka-streams-test-utils` dependencies that we use to enable testing.

Authoring the Processor

5. Have a look at the content of subfolder `src/main/java/streams`. You should find 3 Java files in it:

- ! `CustomMaxAggregatorSupplier.java`

- ! `ConfigProvider.java`

- ! `TopologyProvider.java`

6. Open the class `CustomMaxAggregatorSupplier` and have a look into the code and try to understand what's happening. This is the code we're going to test ultimately. Maybe discuss the code with your peers.
7. Now a Kafka Streams application also needs some configuration. For this purpose we have the `ConfigProvider` class. Once again make sure you understand the code before you proceed.
8. Finally we have a `TopologyProvider` class which defines the topology of the Kafka Streams app that we want to test. And again we invite you to analyze the code and discuss it with your peers if needed.
9. Due to the fact that we use this topology in a test scenario it has some settings that would not be recommended in production. Can you spot them? If yes, discuss how we could improve this class to work well for both scenarios, testing and production run.
10. Now we're ready for the actual test class. Open the class `ProcessorTest.java` located in subfolder `src/test/java/streams`. This file contains the skeleton of the test class.

Let's discuss the code:

- ! We are using the `TopologyTestDriver` class to test the topology. This is a helper class from the `kafka-streams-test-utils` library.

- ! To send records to the test driver we are using the `TestInputTopic` class.

- ! To verify the results of our application we are using the `TestOutputTopic` class.

- ! In the `setup()` method, we're using our two classes `ConfigProvider` and

`TopologyProvider` to get the configuration and topology of our Kafka Streams application.

- ! With the latter two we create a test driver instance that we will use in our tests
- ! The processor is stateful and we pre-populate the state store with a value of `21` for the key `a`.
- ! In the tear down method we simply clean up by closing the test driver instance.

Writing a Test

Now we are ready to write our first test.

You may want to examine the Javadocs for the following classes we will be using:

TopologyTestDriver:

<https://kafka.apache.org/25/javadoc/org/apache/kafka/streams/TopologyTestDriver.html>

TestInputTopic:

<https://kafka.apache.org/25/javadoc/org/apache/kafka/streams/TestInputTopic.html>

TestOutputTopic:

<https://kafka.apache.org/25/javadoc/org/apache/kafka/streams/TestOutputTopic.html>

11. The first test method to the `ProcessorTest` class will test whether reading from the `result-topic` will give the current max value of `21` for the key `a` after inputting a smaller value for the same key. If so, this assures us that the first input flushed to the local state store and the result was produced to the output topic. You can choose to implement this test yourself by researching the `TopologyTestDriver`, `TestInputTopic`, and `TestOutputTopic` classes as well as the `org.hamcrest.MatcherAssert` and `org.hamcrest.CoreMatchers` classes before moving on.

```
@Test
public void shouldFlushStoreForFirstInput() {
    // TODO: add test code here
}
```

So far this is just standard `JUnit` test method.

12. Add code inside the above method to create an input record using the `pipelineInput` method of the `TestInputTopic` class:

```
inputTopic.pipelineInput("a", 1L);
```

13. We can use the `readKeyValue` method of the `TestOutputTopic` class to read the output record generated by the processor and the `assertThat` method to compare the key value with the expected result

```
assertThat(outputTopic.readKeyValue(), equalTo(new KeyValue<>("a", 21L)));
```

14. Finally we make sure that this was the only result that we got as output for the given input:

```
assertThat(outputTopic.isEmpty(), is(true));
```

15. The final test method should look like this:

```
@Test
public void shouldFlushStoreForFirstInput() {
    // TODO: add test code here
    inputTopic.pipelineInput("a", 1L);
    assertThat(outputTopic.readKeyValue(), equalTo(new KeyValue<>("a", 21L)));
    assertThat(outputTopic.isEmpty(), is(true));
}
```

And that's it! We have just authored a complete test.

Running the Test(s) & Displaying the Test Report

16. In the terminal window:

```
$ ./gradlew test
```

The output of this command should look similar to this:

```
BUILD SUCCESSFUL in 4s
3 actionable tasks: 2 executed, 1 up-to-date
```

As we can see, the code compiled and the tests ran without a problem.

17. Navigate to the `simple-test/build/reports/tests/test/index.html` test report in a browser.



Open from file explorer providing the full path `confluent-streams/labs/testing/simple-test/build/reports/tests/test/`

it should look similar to this:

At this point, we should see 1 test that ran successfully.

Adding more Tests

18. Provided is a second test that makes sure that the store value is not updated for smaller input values

```
@Test
public void shouldNotUpdateStoreForSmallerValue() {
    inputTopic.pipelineInput("a", 1L);
    assertEquals(store.get("a"), equalTo(21L));
    assertEquals(outputTopic.readKeyValue(), equalTo(new KeyValue<>("a", 21L)));
    assertTrue(outputTopic.isEmpty(), is(true));
}
```

Discuss with your peers what this test does.

19. Add more tests and repeat the `./gradlew test` command to see new results.

a. Add a test that asserts that the store value is updated for a larger input value.



If you get stuck, please have a look in the solutions folder where you will find the complete sample solution.

- b. Add a test that asserts a new store value is generated if adding a value with a new key "b". Also make sure the existing store value for "a" is unchanged.
- c. Write a test that verifies that the processor punctuates if the event time advances.
- d. Write a test that verifies that the processor punctuates if the wall clock time advances.



Use the function `testDriver.advanceWallClockTime()` for this.

Conclusion

In this sample we have shown how to test a simple Kafka Streams application that uses the Processor API.

b. Integration Tests using Embedded Kafka

In this exercise, we are going a step further to use a full blown embedded Kafka single node cluster to test our Kafka Streams application. The application is a simple word count streaming app. As in the previous test, we do not need to run our Docker container Kafka cluster.

Prerequisites

Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)

Creating the Artifacts to Test

1. Navigate to folder `~/confluent-streams/labs/testing/wordcount-test`, and launch VS Code:

```
$ cd ~/confluent-streams/labs/testing/wordcount-test
$ code .
```

2. Open the file `build.gradle` and analyze its content. Notice the many dependencies that we have to add to enable this kind of integration testing.

!

we have a few dependencies that use the test version of the respective `JAR` files. This can be achieved by adding `classifier: 'test'` or `classifier: 'tests'` to the dependency.

3. Open the file `TopologyProvider.java` class, located in subfolder `src/main/java/streams`, that defines the Kafka Streams topology that we will test. Analyze the code. Make sure you understand it. If not discuss it with your peers.

Writing the Integration Test

Now it is time to write the actual test that is leveraging the testbed that we have just prepared.

4. Locate the file `TopologyProviderTest.java` in subfolder `src/test/java/streams` and open it. It will host our test code.

Let's analyze the code a bit:

- ! In the `@BeforeClass` we're initializing and running an embedded single node Kafka cluster. We are preparing the cluster by creating the input and the output topic.
- ! The method `shouldCountWords` contains our first test. We have clearly documented the steps:
 1. define the configuration of the streaming app to use during the test
 2. get the topology that we want to test
 3. initialize and start the streaming application
 4. produce some data for the input topic
 5. finally verify that the produced data in the output topic corresponds to the expected values
- 5. Inspect the method `getStreamsConfiguration`. Notice that we take the information about the bootstrap server(s) from our `CLUSTER` variable which represents our single node embedded Kafka cluster. We also provide a path to where the state store should store its information (`STATE_DIR_CONFIG`).
- 6. Inspect the method `produceInputData`. Notice that we're configuring a producer and are using it to feed the data into the input topic.
- 7. The final step is to write code that validates the results. Inspect the `verifyOutputData` method. You can challenge yourself to complete the method, or you can proceed to the next step.
- 8. Make it so the `verifyOutputData` method has these contents:


```

List<KeyValue<String, Long>> expectedWordCounts = Arrays.asList(
    new KeyValue<>("hello", 1L),
    new KeyValue<>("all", 1L),
    new KeyValue<>("streams", 2L),
    new KeyValue<>("lead", 1L),
    new KeyValue<>("to", 1L),
    new KeyValue<>("join", 1L),
    new KeyValue<>("kafka", 3L),
    new KeyValue<>("summit", 1L)
);

Properties consumerConfig = new Properties();
consumerConfig.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, CLUSTER
    .bootstrapServers());
consumerConfig.put(ConsumerConfig.GROUP_ID_CONFIG,
    "wordcount-lambda-integration-test-standard-consumer");
consumerConfig.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG,
    "earliest");
consumerConfig.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
    StringDeserializer.class);
consumerConfig.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
    LongDeserializer.class);

List<KeyValue<String, Long>> actualWordCounts =
    IntegrationTestUtils.waitUntilMinKeyValueRecordsReceived(
        consumerConfig, outputTopic, expectedWordCounts.size());
assertThat(actualWordCounts, containsInAnyOrder(expectedWordCounts
    .toArray()));

```

We define an array with the expected data, define a consumer configuration, and then compare the produced data to the expected data.

Running the Test(s)

9. In the terminal window:

```
$ ./gradlew test
```

You should see something like this (shortened):

```
...
> Task :compileJava
> Task :processResources NO-SOURCE
> Task :classes
> Task :compileTestJava
> Task :processTestResources NO-SOURCE
> Task :testClasses
> Task :test

BUILD SUCCESSFUL in 46s
3 actionable tasks: 3 executed
...
```

10. Navigate to the `wordcount-test/build/reports/tests/test/index.html` folder and open the `index.html` test report in a browser.

Cleanup

There is no special cleanup needed.

Conclusion

In this exercise we have written an integration test for a Kafka Streams app topology. The test bed uses an embedded version of Kafka and Zookeeper.

STOP HERE. THIS IS THE END OF THE EXERCISE.

c. Getting Metrics from a Kafka Streams Application

The purpose of this exercise is to learn how to expose metrics in a Kafka Streams application. This particular application sends metrics to standard output every 10 seconds. In practice, these metrics can be exposed to external sources for aggregation and analysis, as we will see in the exercises that follow.

Preparing the application

In this exercise we're going to use the word count exercise from a previous lab.

1. Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)
2. Navigate to the folder for this lab:

```
$ cd ~/confluent-streams/labs/monitoring
```

3. Run the Kafka cluster:

```
$ docker-compose up -d zookeeper kafka control-center
```

4. Create the two topics called `lines-topic` and `word-count-topic` in Kafka:

```
$ kafka-topics \
E --create \
E --bootstrap-server kafka:9092 \
E --replication-factor 1 \
E --partitions 1 \
E --topic lines-topic

$ kafka-topics \
E --create \
E --bootstrap-server kafka:9092 \
E --replication-factor 1 \
E --partitions 1 \
E --topic word-count-topic
```

Building & Running the Application

5. In a terminal window navigate to the `word-count` folder, and launch VS Code:

```
$ cd ~/confluent-streams/labs/monitoring/word-count
$ code .
```

6. Open the file `build.gradle` in folder `word-count` and analyze its content. It should be quite familiar by now.
7. Notice the four Java files in subfolder `src/main/java/streams`:

<code>ConfigProvider.java</code>	Defines the configuration for the streams application
<code>MetricsReporter.java</code>	Defines how the list of metrics is output every 10 seconds by the app
<code>TopologyProvider.java</code>	Defines the topology of the stream application
<code>WordCountSample.java</code>	Main class of the application. This class makes use of the 3 following classes

Analyze their code. Specifically note the use of the `MetricsReporter` class. Make sure you understand what's going on in the code.

8. Use `Run ! Start Debugging` in VS Code or `./gradlew run` in the terminal to run your code.

The application will print out the list of metrics to the terminal every 10 seconds. It should look similar to this (shortened for readability):

```
--- Application Metrics ---
MetricName [name=count, group=kafka-metrics-count, description=total
number of registered metrics, tags={client-id=wordCount-3e02d8d9-
490b-492a-9bf0-cf85629e7fcf-StreamThread-1-producer}], 81.0
MetricName [name=io-time-ns-avg, group=producer-metrics,
description=The average length of time for I/O per select call in
nanoseconds., tags={client-id=wordCount-3e02d8d9-490b-492a-9bf0-
cf85629e7fcf-StreamThread-1-producer}], 612250.0
```

Producing Input Data

To see how the metrics change, we can produce some input data that the application will process. We use the `kafka-console-producer` tool for this job.

9. Return to the terminal window and create a list of input sentences that will be randomly produced to the input topic:

```
$ INPUTS=('Kafka powers the Confluent streaming platform' \
Ê 'All Streams come from Kafka' \
Ê 'Streams will all flow to Kafka' \
Ê 'Follow the streams to Kafka Summit' \
Ê 'Check out Confluent Cloud')
```

10. Run `Kafkacat` to produce a steady flow of sentences to the input topic:

```
$ while true; do
Ê MESSAGE=${INPUTS[${RANDOM} % ${#INPUTS[@]}]}
Ê echo ${MESSAGE} | kafkacat -P \
Ê -b kafka:9092 \
Ê -t lines-topic
Ê sleep 0.1
Ê done
```

11. Observe the metric values printed by the Word Count application.



Use `grep` to track a metric. For example, run: `./gradlew run | grep name=poll-records-avg` to track the metric `poll-records-avg` and then start/stop the `kafkacat` producer of Step 10.

12. In another terminal window run `kafka-console-consumer` to report the output of our

sample app:

```
$ kafka-console-consumer --bootstrap-server kafka: 9092 \  
E --topic word-count-topic \  
E --property print.key=true \  
E --value-deserializer  
org.apache.kafka.common.serialization.LongDeserializer
```

Cleanup

13. Quit the Producer with **Ctrl +C**.
14. Quit the Consumer with **Ctrl +C**.
15. Quit our sample app with Run ! Stop Debugging.
16. Shut down your Kafka cluster with the **docker-compose down -v** command.

d. Using JConsole to monitor a Streams App

In this exercise, we will use **JConsole** to monitor the various metrics a simple Kafka Streams application exposes.

We will be using an application we created in an earlier lab that reads data from a topic whose keys are integers and whose values are sentence strings. The input values are transformed to lower-case and output to a new topic.

Prerequisites

Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)

Preparing the Lab

1. Navigate to the module folder:

```
$ cd ~/confluent-streams/labs/monitoring
```

2. Run the Kafka cluster:

```
$ docker-compose up -d zookeeper kafka
```

3. Create an input topic called **lines-topic** in Kafka:

```
$ kafka-topics \
E   --create \
E   --bootstrap-server kafka:9092 \
E   --replication-factor 1 \
E   --partitions 1 \
E   --topic lines-topic
```

4. Create the output topic called **lines-lower-topic** in Kafka:

```
$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic lines-lower-topic
```

5. In a terminal window navigate to the project folder `jmx-sample` and build the artifact:

```
$ cd ~/confluent-streams/labs/monitoring/jmx-sample
$ ./gradlew build
```

6. Examine `build.gradle` and observe it has been configured to expose JMX metrics on port `4444` that we can then use to attach `JConsole`:

```
$ cat build.gradle
...
applicationDefaultJvmArgs = [
  "-Dcom.sun.management.jmxremote",
  "-Dcom.sun.management.jmxremote.authenticate=false",
  "-Dcom.sun.management.jmxremote.ssl=false",
  "-Djava.rmi.server.hostname=127.0.0.1",
  "-Dcom.sun.management.jmxremote.rmi.port=4444",
  "-Dcom.sun.management.jmxremote.port=4444" ]
...
```

7. Still in the project folder `jmx-sample` run your Kafka Streams application:

```
$ ./gradlew run
```

8. Observe the JMX metrics of the Kafka Streams application:

- a. Open a new Terminal window and open a `jconsole` connection to port `4444` which is the JMX port for the Kafka Streams application.

```
$ jconsole localhost:4444 &
```

- b. Select Insecure connection when asked



- c. Navigate to the MBeans tab
- d. Explore the node under kafka-streams as indicated in the below image:

Initially some of the numbers, such as `process total` will be zero.

9. Open a new terminal window and navigate to the `monitoring` folder:

```
$ cd ~/confluent-streams/abs/monitoring
```

10. Now create some data that will be consumed and processed by the Kafka Streams application:

```
$ cat << EOF | kafka-console-producer \  
E  --bootstrap-server kafka:9092 \  
E  --property "parse.key=true" \  
E  --property "key.separator=: " \  
E  --topic lines-topic \  
1: "Kafka powers the Confluent Streaming Platform" \  
2: "Events are stored in Kafka" \  
3: "Confluent contributes to Kafka" \  
EOF
```

11. Observe how the values of the metrics in JConsole change (you will need to click the Refresh button):

12. Optional: add more data to the topic and monitor the attributes and how the values change.

Cleanup

13. Quit JConsole.
14. Stop the sample Kafka Streams application by pressing `Ctrl +C`.
15. Shut down your Kafka cluster with the `docker-compose down -v` command.

Conclusion

We have used Java code to retrieve the metrics directly from the `KafkaStreams` object and also configured the Kafka Streams sample application to expose JMX data that we then explored using `JConsole`.

e. Monitoring a Kafka Streams App in Confluent Control Center

In this exercise we're going to reuse the word count processor API example application from a previous exercise and extend it so that it can be monitored in Confluent Control Center.

Prerequisites

1. Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)
2. In a terminal window navigate to the `processor-sample` folder:

```
$ cd ~/confluent-streams/labs/monitoring/processor-sample
```

3. Open this application's root directory in VS Code.

```
$ code .
```

4. Open the `build.gradle` file and observe `monitoring-interceptors` in the list of dependencies:

```
compile group: "io.confluent", name: "monitoring-interceptors",  
version: "6.0.0"
```

This dependency contains the interceptors classes that we will use to configure our application for monitoring via Confluent Control Center.

5. Open the file `ProcessorSample.java` (in folder `src/main/java/streams`) which contains the `main` function of the sample application. Notice these lines that have been added to the `getConfig` function:

```
settings.put(StreamsConfig.producerPrefix(ProducerConfig.INTERCEPTOR_CLASSES_CONFIG),
    "io.confluent.monitoring.clients.interceptor.MonitoringProducerInterceptor");
settings.put(StreamsConfig.consumerPrefix(ConsumerConfig.INTERCEPTOR_CLASSES_CONFIG),
    "io.confluent.monitoring.clients.interceptor.MonitoringConsumerInterceptor");
```



Please note the prefix used for both consumer and producer interceptors.

With interceptors configured, we can monitor our app through Confluent Control Center.

6. Start the Kafka cluster and the Confluent Control Center server with:

```
$ docker-compose up -d zookeeper kafka control-center
```

As usual, you are encouraged to inspect the content of the `docker-compose` file in the `~confluent-streams` directory to make sure you understand all the settings and discuss it with your peers. Wait a couple of minutes until the cluster is initialized. Open Confluent Control Center at <http://localhost:9021> and wait until it displays the system health.

7. Create the input and output topics called `lines-topic` and `word-count-topic` respectively. For the moment let's give them each 1 partition and replication factor 1:

```
$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic lines-topic

$ kafka-topics \
  --create \
  --bootstrap-server kafka:9092 \
  --replication-factor 1 \
  --partitions 1 \
  --topic word-count-topic
```

8. Use Run ! Start Debugging in VS Code or `./gradlew run` in the terminal to run your

streams app.



Ignore the WARNINGS.

Producing Data

9. Run the `kafka-console-producer` that we will use to feed some data to the topic `lines-topic`:

```
$ NS=io.confluent.monitoring.clients.interceptor && \
kafka-console-producer \
  --bootstrap-server kafka:9092 \
  --topic lines-topic \
  --producer-property \
  "interceptor.classes=${NS}.MonitoringProducerInterceptor"
```

10. Open another terminal tab and run the `kafka-console-consumer` for the `word-count-topic` topic:

```
$ NS=io.confluent.monitoring.clients.interceptor && \
kafka-console-consumer \
  --group word-count-consumer \
  --bootstrap-server kafka:9092 \
  --topic word-count-topic \
  --property print.key=true \
  --consumer-property \
  "interceptor.classes=${NS}.MonitoringConsumerInterceptor"
```

Note the use of the group name `word-count-consumer` and the `MonitoringConsumerInterceptor` class to enable monitoring of the consumer in Control Center.

11. In the terminal window where the producer runs, enter a few lines of text such as:

```
Kafka is powering the Confluent streaming platform
Streaming in real-time is more and more important
Our company will invest in real-time streaming
For this we need to know loads about Kafka and Kafka Streams
ksqlDB is a simpler alternative to Kafka Streams
Everybody loves ksqlDB since no programming is required
```

and observe the output in the terminal window where the consumer runs.



Add more data at will so that there is some activity ongoing.

12. In Confluent Control Center (<http://localhost:9021>) navigate to Consumers. You should see something like this:

Click on our Kafka Streams (processor-sample-v0.1.0) application to investigate its metrics.

Cleanup

13. To stop the producer hit **Ctrl +C**.
14. To stop the consumer hit **Ctrl +C**.
15. To stop the Kafka Streams application with Run ! Stop Debugging.
16. To stop the Kafka Cluster and delete the volumes execute this command:

```
$ cd ~/confluent-streams/abs/monitoring/processor-sample
$ docker-compose down -v
```


STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 11 Security

a. Securing a Kafka Streams Application

In this exercise, we will interact with a secure Kafka cluster in several ways. The cluster is configured to use SASL-PLAIN for authentication and SSL for transport encryption. A certificate creation script is used to create keystores and truststores for clients and brokers so that they can authenticate with each other. There is a client `.properties` that will be used to create topics, to produce to an input topic, and to consume from an output topic. There is another `.properties` file to configure the security for a Kafka Streams application. In practice, SSL certificates must be carefully managed, credentials should be created separately for different clients, and permissions should be applied to secure credentials. These areas are beyond the scope of the course. Rather, the purpose of the exercise is to provide hands-on experience with the configurations necessary to connect a Kafka Streams application to a secured cluster.

Preparing the Project

Please make sure you have prepared your lab environment as described here: ! [Lab Environment](#)

#

If you have a cluster running from an earlier lab, shut it down now. This lab depends on the settings in the local `docker-compose.yml` file. You can check if you have a cluster running with `docker-compose ps` and you can shutdown a running cluster with `docker-compose down -v`.

Preparing the Certificates

First we need to generate all the necessary certificates and credentials that are used to create a secure (single node) Kafka cluster.

1. In a terminal window, navigate to the folder `security/secure-app/scripts/security`:

```
$ cd ~/confluent-streams/labs/security/secure-app/scripts/security
```

2. Take a few minutes to inspect the contents of the files in this directory and discuss them with your peers.
3. Run the `certs-create.sh` script to generate the necessary certs and credential files:

```
$ ./certs-create.sh
```

Running the Cluster

4. In a terminal window navigate to folder `security/secure-app`, and launch VS Code:

```
$ cd ~/confluent-streams/labs/security/secure-app  
$ code .
```

5. From this folder open the file `docker-compose.yml` and analyze it. Specifically focus on the settings of the security relevant environment variables. Discuss the content with your peers if you don't fully understand it.
6. Start the cluster:

```
$ docker-compose up -d zookeeper kafka control-center
```

and wait a couple of minutes until the cluster is initialized. You may observe the progress by using:

```
$ docker-compose ps
```

or following the logs, e.g.:

```
$ docker-compose logs -f kafka
```

Press `Ctrl +C` to stop following the log.

Creating the Application

7. In the `secure-app` folder, locate the file `build.gradle` and analyze its content. There should be no surprises at this point.
8. Open the file `SecureAppSample.java` located in the subfolder `src/main/java/streams` and inspect its contents. As you can see, we're using a class `ConfigProvider` to get us the configuration for the Kafka Streams app and a class `TopologyProvider` to get us the topology for the application. Other than that all is well known and familiar code by now.
9. Now let's have a look at the file `ConfigProvider.java`. Notice that this code gets the configuration settings from a `.properties` file, which is preferable to hard-coding properties into our application.
10. In the project sub-folder `scripts/security`, open the file called `secureapp-sample.properties` analyze its content.
 - ! Note that we're using SASL to authenticate our application with the secured Kafka cluster.
 - ! The selection of the protocol (SASL_SSL) also has an implication on which port we're using for the communication with the brokers; `9091` in this case. The brokers are also listening to `SSL` on port `11091`, if we wanted to do mutual SSL instead of SASL + SSL. In that case, we would also have to configure `ssl.keystore` location and password on the clients so that Brokers could authenticate them.
 - ! We need to define the trust store and its password as well as the login module and the credentials.
 - ! We're also securing the monitoring interceptors so that the app can be monitored from Confluent Control Center.
11. Copy the security files to the location in `/etc/kafka/secrets/` - remember our user ID is training and the password which will be requested is also training:

```
$ sudo mkdir -p /etc/kafka/secrets/  
$ sudo cp scripts/security/* /etc/kafka/secrets/
```

12. Finally have a look at the file `TopologyProvider.java` which should have this content:

```

package streams;

import org.apache.kafka.streams.Topology;
import org.apache.kafka.streams.StreamsBuilder;

public class TopologyProvider {
    public Topology getTopology() {
        final StreamsBuilder builder = new StreamsBuilder();
        builder.stream("secure-input").to("secure-output");
        return builder.build();
    }
}

```

!

Since this sample is all about securing a Kafka Streams application we have chosen a very simple Topology, it's basically a **NO-OP** topology. We read input from topic **secure-input** and directly and unmodified output all the messages to the topic **secure-output**.

13. Before we can run our Kafka Streams sample application we need to first manually create the topics **secure-input** and **secure-output** since we have configured the Kafka cluster to disable auto-create of topics (see setting in the **docker-compose.yml** file). The **AdminClient** that creates topics must also authenticate with the Kafka cluster, so we must use the **--command-config** option to point to the **.properties** file that contains the security configurations.

```

$ kafka-topics \
  --create \
  --bootstrap-server kafka:9091 \
  --replication-factor 1 \
  --partitions 3 \
  --topic secure-input \
  --command-config /etc/kafka/secrets/client_security.properties

$ kafka-topics \
  --create \
  --bootstrap-server kafka:9091 \
  --replication-factor 1 \
  --partitions 3 \
  --topic secure-output \
  --command-config /etc/kafka/secrets/client_security.properties

```

#

You will see some warnings such as

```
2018-10-15 18:12:39 WARN ProducerConfig:287 - The
configuration
'confluent.monitoring.interceptor.security.protocol' was
supplied but isn't a known config.
```

You can safely ignore these warnings. The properties are necessary monitoring producers and consumers in Confluent Control Center, but are not recognized by the core open source Apache Kafka `ProducerConfig` and the `ConsumerConfig` classes.

14. Now we're ready to run the app. Use Run ! Start Debugging to run your Kafka Streams application.

!

Ignore the WARNINGS.

Creating Input Data

To see the sample streams application working we need to generate some data.

15. In the `~/confluent-streams/labs/security/secure-app/scripts/security` folder, there is a file called `client_security.properties` with the following content:

```

bootstrap.servers=kafka: 9091
security.protocol =SASL_SSL
ssl.truststore.location=/etc/kafka/secrets/kafka.client.truststore.jks
ssl.truststore.password=confluent
sasl.mechanism=PLAIN
sasl.jaas.config=org.apache.kafka.common.security.plain.PlainLoginModule required \
    username=\"client\" \
    password=\"client-secret\";

# authenticate the monitor interceptor with Kafka.
confluent.monitoring.interceptor.bootstrap.servers=kafka: 9091
confluent.monitoring.interceptor.security.protocol =SASL_SSL
confluent.monitoring.interceptor.ssl.truststore.location=/etc/kafka/secrets/kafka.client.truststore.jks
confluent.monitoring.interceptor.ssl.truststore.password=confluent
confluent.monitoring.interceptor.sasl.mechanism=PLAIN
confluent.monitoring.interceptor.sasl.jaas.config=\
    org.apache.kafka.common.security.plain.PlainLoginModule required \
    username=\"client\" \
    password=\"client-secret\";

```

In an earlier step, we copied this file, along with others, to `/etc/kafka/secrets/`. This file will be used by the client tools that we're going to use in a moment to authenticate themselves, as it was earlier when we created the secure-input and secure-output topics.

16. Open a new terminal window, navigate to the `secure-app` folder:

```
$ cd ~/confluent-streams/labs/security/secure-app
```

17. To run the `kafka-console-producer` that we will use to feed some data to the topic `secure-input` use this command:

```

$ NS=io.confluent.monitoring.clients.interceptor && \
kafka-console-producer \
    --bootstrap-server kafka: 9091 \
    --topic secure-input \
    --producer.config /etc/kafka/secrets/client_security.properties \
    --producer-property \
    interceptor.classes="${NS}.MonitoringProducerInterceptor"

```

!

- ! We're configuring the producer for monitoring via Control Center through the parameter `--producer-property`
- ! We're passing the properties file with the security settings through the parameter `--producer.config` to the producer.

18. Open another terminal window, navigate to the `secure-app` folder:

```
$ cd ~/confluent-streams/labs/security/secure-app
```

19. To list the data produced by our sample application to the topic `secure-output` use this command:

```
$ NS=io.confluent.monitoring.clients.interceptor && \
kafka-console-consumer \
  --group secure-console-consumer \
  --bootstrap-server kafka:9091 \
  --topic secure-output \
  --from-beginning \
  --consumer.config /etc/kafka/secrets/client_security.properties \
  --consumer-property \
  interceptor.classes="${NS}.MonitoringConsumerInterceptor"
```

!

- ! We're configuring the consumer for monitoring via Control Center through the parameter `--consumer-property`
- ! We're passing the properties file with the security settings through the parameter `--consumer.config` to the consumer

20. In the terminal where the producer is running enter a few lines such as:

```
Kafka is powering the Confluent streaming platform
Real time streaming is exceedingly important
Confluent offers support for Kafka Streaming and KSQL DB
In my company we build Kafka Streams applications
```

21. Observe how output is generated in the terminal window where the `kafka-console-consumer` is running. Specifically notice that the order of the output relative to the input might be changed (if you enter values fast enough) due to the fact that we have 3

partitions per topic and ordering is only guaranteed within a partition but not globally per topic!

22. Optional: Add a lot more data using the `kafka-console-producer`. For example:

```
$ NS=io.confluent.monitoring.clients.interceptor
$ for i in {1..100}; do
$   sleep 1
$   seq 1000 | kafka-console-producer \
$     --bootstrap-server kafka:9091 \
$     --topic secure-input \
$     --producer.config /etc/kafka/secrets/client_security.properties \
$     --producer-property \
$     interceptor.classes="${NS}.MonitoringProducerInterceptor"
done
```

23. Open Confluent Control Center at <http://localhost:9021> and go to Consumers !
secureapp-sample and you should see something like this:

Cleanup

24. Quit both the producer and consumer by pressing `Ctrl +C`.
25. Quit the sample Kafka Streams application with Run ! Stop Debugging.
26. Shutdown the Kafka cluster:

```
$ docker-compose down -v
```

Conclusion

In this example we have shown how to run a secure Kafka cluster and then build a Kafka Streams application that integrates with this cluster using SASL for authentication and SSL for encryption. The application logic was trivial yet that is not the point of this example. The important fact is the integration with the Kafka cluster security settings.

STOP HERE. THIS IS THE END OF THE EXERCISE.

Appendix A: Running All Labs with Docker

Running Labs in Docker for Desktop

If you have installed Docker for Desktop on your Mac or Windows 10 Pro machine you are able to complete the course by building and running your applications from the command line.

- ¥ Increase the memory available to Docker Desktop to a minimum of 6 GiB. See the advanced settings for [Docker Desktop for Mac](#), and [Docker Desktop for Windows](#).
- ¥ Follow the instructions at ! [The Lab Environment & Sample Solutions](#) to `git clone` the source code, in each exercise follow the instructions to launch the cluster containers with `docker-compose` on your host machine. The exercise source code will now be on your host machine where you can edit the source code with any editor.
- ¥ Begin the exercises by first opening a bash shell on the tools container. All the command line instructions will work from the tools container. This container has been preconfigured with all of the tools you use in the exercises, e.g. `kafka-topics` and `python`.

```
$ docker-compose exec tools bash
bash-4.4#
```

- ¥ At the time of writing, the Python prerequisite `confluent_kafka` won't install with the version of Python 3 available in the tools container. You should use `pip` and `python` rather than `pip3` and `python3`.
- ¥ At the time of writing, Maven is not installed in the tools container, so the optional maven exercise would require a hefty 300 MB download with `apt-get install maven`. It might be best to skip this optional exercise and simply use it as reference if you use Maven in your day-to-day work.
- ¥ Anywhere you are instructed to open additional terminal windows you can `exec` additional bash shells on the tools container with the same command as above on your host machine.

¥ Any subsequent `docker` or `docker-compose` instructions should be run on your host machine.

Running the Exercise Applications

From the `tools` container you can use command line alternatives to the VS Code steps used in the instructions. Complete the exercise code with an editor on your host machine, then use the following command line instructions to build and run the applications from the exercise directory.

¥ For Java applications: `./gradlew run`

¥ For Python applications: `python main.py`

Where you are instructed to use Run ! Stop Debugging in VS code, use `Ctrl +C` to end the running exercise.

Our `docker-compose.yml` file sets the working directory inside the tools container to the `~/confluent-streams directory`. Be sure to change the working directory to each exercise directory as stated in each exercise instructions.

```
$ docker-compose exec tools bash
bash-4.4# pwd
/root/confluent-streams
```

To build and run the Anatomy of a Kafka Streams App exercise. First make the source code updates to `~/confluent-streams/labs/streams-writing/gradle-sample/src/main/java/streams/MapSample.java` on your host machine. Then enter into the bash shell on your tools container:

```
bash-4.4# cd labs/streams-writing/gradle-sample
bash-4.4# ./gradlew run
```

In the Monitoring Kafka Streams Applications exercise, you must use `jconsole` on port 4444 on the host to view JMX metrics. If `jconsole` is not already installed on your host system, install using the password `training`:

```
$ sudo apt install -y openjdk-11-jdk
```

Run a new tools container with a port mapping to expose JMX metrics to the host:

```
$ docker-compose run -p 4444:4444 tools  
bash-4.4# cd jmx-sample && ./gradlew run
```

The application will now expose metrics to port 4444 in the container, which is mapped to port 4444 on the host. Running **jconsole** on the host on port 4444 will now pick up the metrics exposed by the application.